

# Multi-Core Programming

*For Academia*

**Student Workbook**

---

*Intel® Software College*





# Multi-Core Programming

*For Academia*

Student Workbook

---

*Intel® Software College*



The information contained in this document is provided for informational purposes only and represents the current view of Intel Corporation ("Intel") and its contributors ("Contributors") on, as of the date of publication. Intel and the Contributors make no commitment to update the information contained in this document, and Intel reserves the right to make changes at any time, without notice.

DISCLAIMER. THIS DOCUMENT, IS PROVIDED "AS IS." NEITHER INTEL, NOR THE CONTRIBUTORS MAKE ANY REPRESENTATIONS OF ANY KIND WITH RESPECT TO PRODUCTS REFERENCED HEREIN, WHETHER SUCH PRODUCTS ARE THOSE OF INTEL, THE CONTRIBUTORS, OR THIRD PARTIES. INTEL, AND ITS CONTRIBUTORS EXPRESSLY DISCLAIM ANY AND ALL WARRANTIES, IMPLIED OR EXPRESS, INCLUDING WITHOUT LIMITATION, ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR ANY PARTICULAR PURPOSE, NON-INFRINGEMENT, AND ANY WARRANTY ARISING OUT OF THE INFORMATION CONTAINED HEREIN, INCLUDING WITHOUT LIMITATION, ANY PRODUCTS, SPECIFICATIONS, OR OTHER MATERIALS REFERENCED HEREIN. INTEL, AND ITS CONTRIBUTORS DO NOT WARRANT THAT THIS DOCUMENT IS FREE FROM ERRORS, OR THAT ANY PRODUCTS OR OTHER TECHNOLOGY DEVELOPED IN CONFORMANCE WITH THIS DOCUMENT WILL PERFORM IN THE INTENDED MANNER, OR WILL BE FREE FROM INFRINGEMENT OF THIRD PARTY PROPRIETARY RIGHTS, AND INTEL, AND ITS CONTRIBUTORS DISCLAIM ALL LIABILITY THEREFOR.

INTEL, AND ITS CONTRIBUTORS DO NOT WARRANT THAT ANY PRODUCT REFERENCED HEREIN OR ANY PRODUCT OR TECHNOLOGY DEVELOPED IN RELIANCE UPON THIS DOCUMENT, IN WHOLE OR IN PART, WILL BE SUFFICIENT, ACCURATE, RELIABLE, COMPLETE, FREE FROM DEFECTS OR SAFE FOR ITS INTENDED PURPOSE, AND HEREBY DISCLAIM ALL LIABILITIES THEREFOR. ANY PERSON MAKING, USING OR SELLING SUCH PRODUCT OR TECHNOLOGY DOES SO AT HIS OR HER OWN RISK.

Licenses may be required. Intel, its contributors and others may have patents or pending patent applications, trademarks, copyrights or other intellectual proprietary rights covering subject matter contained or described in this document. No license, express, implied, by estoppels or otherwise, to any intellectual property rights of Intel or any other party is granted herein. It is your responsibility to seek licenses for such intellectual property rights from Intel and others where appropriate.

Limited License Grant. Intel hereby grants you a limited copyright license to copy this document for your use and internal distribution only. You may not distribute this document externally, in whole or in part, to any other person or entity.

LIMITED LIABILITY. IN NO EVENT SHALL INTEL, OR ITS CONTRIBUTORS HAVE ANY LIABILITY TO YOU OR TO ANY OTHER THIRD PARTY, FOR ANY LOST PROFITS, LOST DATA, LOSS OF USE OR COSTS OF PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES, OR FOR ANY DIRECT, INDIRECT, SPECIAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF YOUR USE OF THIS DOCUMENT OR RELIANCE UPON THE INFORMATION CONTAINED HEREIN, UNDER ANY CAUSE OF ACTION OR THEORY OF LIABILITY, AND IRRESPECTIVE OF WHETHER INTEL, OR ANY CONTRIBUTOR HAS ADVANCE NOTICE OF THE POSSIBILITY OF SUCH DAMAGES. THESE LIMITATIONS SHALL APPLY NOTWITHSTANDING THE FAILURE OF THE ESSENTIAL PURPOSE OF ANY LIMITED REMEDY.

Intel and Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.

\*Other names and brands may be claimed as the property of others.

Copyright © 2007, Intel Corporation. All Rights Reserved.

## Contents

---

<b>Lab. 1: Intel® VTune™ Performance Analyzer .....</b>	<b>9</b>
<b>1: Finding Hotspots .....</b>	<b>11</b>
Collect Sampling Data for the Clockticks Event .....	11
Review Questions .....	12
<b>2: Sampling Over Time .....</b>	<b>13</b>
Create a Sampling Activity .....	13
<b>3: Using Call Graph .....</b>	<b>14</b>
Collect call graph data .....	14
<b>4: Using the Windows Command Line Interface .....</b>	<b>15</b>
Collect Samples based on Clockticks Event using Command Line Interface .....	15
View the Profiling Data for gzip .....	15
Pack the Data and View it in the GUI .....	15
<b>Lab. 2: Intel® Math Kernel Library .....</b>	<b>17</b>
<b>1: Matrix Multiply Sample .....</b>	<b>19</b>
<b>2: Monte Carlo Calculation of Pi .....</b>	<b>20</b>
<b>Lab. 3: Programming with Windows* Threads .....</b>	<b>21</b>
<b>1: Starting with HelloThreads .....</b>	<b>23</b>
Build and Run HelloThreads Program .....	23
Review Questions .....	26
<b>2: Approximating Pi with Numerical Integration .....</b>	<b>27</b>
Build and Run the Serial Program .....	27
Correct Errors and Validate Results .....	27
Review Questions .....	28
<b>3: Using Events .....</b>	<b>29</b>
Build and Run Original Threaded Program .....	29
Modify Original Threaded Program to Use Events .....	29
<b>4: Using Semaphores .....</b>	<b>30</b>
Build and Run Serial Program .....	30
Build and Run Threaded Program .....	30
<b>Lab. 4: Programming with OpenMP Threads .....</b>	<b>31</b>
<b>1: Starting with Hello Worlds .....</b>	<b>33</b>
Initial Compile .....	33
Add OpenMP Directives .....	33
OpenMP Compile .....	33
<b>2: Computing Pi with Numerical Integration .....</b>	<b>35</b>
Initial Compile .....	35
Add OpenMP Directives .....	35
OpenMP Compile and Run .....	36
<b>3: Monte Carlo Pi .....</b>	<b>37</b>


Initial Compile.....	37
Add OpenMP Directives.....	37
OpenMP Compile .....	38
<b>Review Questions .....</b>	<b>39</b>
<b>Lab. 5: Correcting Threading Errors with Intel® Thread Checker for Explicit Threads .....</b>	<b>41</b>
<b>1A: Find Prospective Data Races .....</b>	<b>43</b>
Build and Run Potential Serial Program .....	43
Build and Run Potential Threaded Program .....	43
<b>1B: Resolve Data Races .....</b>	<b>45</b>
Resolve the Problems .....	45
<b>2: Identifying Deadlock.....</b>	<b>46</b>
Build and Run the Program .....	46
Run Application within Intel® Thread Checker .....	48
Correct Errors and Validate Results .....	49
Review Questions .....	50
<b>3: Testing Libraries for Thread Safety .....</b>	<b>51</b>
Setting-up and Compilation for Thread Safety Testing.....	51
Modifying Binary Instrumentation Levels.....	52
Review Questions .....	53
<b>Lab. 6: Tuning Threaded Code with Intel® Thread Profiler for Explicit Threads .....</b>	<b>55</b>
<b>1A: Getting Started with Thread Profiler .....</b>	<b>57</b>
Review Questions .....	58
<b>1B: Analyzing an Application .....</b>	<b>59</b>
<b>2: Finding Load Balance Issues .....</b>	<b>61</b>
Build and Run Potential Threaded Program .....	61
Evaluate Thread Profiler Results to Diagnose Problem .....	61
Fixing the Performance Issue .....	62
<b>3: Finding Synchronization Contention Issues.....</b>	<b>64</b>
Build & Run Threaded Numerical Integration Program .....	64
Evaluate Results to Diagnose Problem .....	64
Using the Timeline View for Analysis.....	67
Fixing the Performance Issue .....	68
Comparing Performance Runs.....	69
<b>Review Questions .....</b>	<b>70</b>
<b>Lab. 7: Threaded Programming Methodology .....</b>	<b>71</b>
<b>1: Compile and Run Serial Code.....</b>	<b>73</b>
Build and Run Serial Program .....	73
<b>2: Analyze Serial Code .....</b>	<b>74</b>
Get Baseline Timing .....	74
Run VTune Performance Analyzer on Application .....	74
<b>3: Run OpenMP Code.....</b>	<b>75</b>

## Contents

Build and Run Threaded Program .....	75
<b>4: Use Thread Checker to Check for Threading Errors .....</b>	<b>76</b>
Use Thread Checker for Analysis.....	76
<b>5: Correct Threading Errors .....</b>	<b>77</b>
Correct Errors and Validate Results.....	77
Build and Run Threaded Program .....	77
<b>6: Find Threading Performance Issues.....</b>	<b>78</b>
Instrument and Run Threaded Program .....	78
Use Thread Profiler for Analysis .....	78
<b>7: Reduce the Number of printf Calls .....</b>	<b>79</b>
Modify Code to Control Number of Progress Updates .....	79
Update Serial Code.....	80
<b>8: Modify Critical Regions .....</b>	<b>81</b>
Use InterlockedIncrement and Local Variables.....	81
<b>9: Fix Load Imbalance .....</b>	<b>83</b>
Use OpenMP schedule Clause .....	83
<b>Lab. 8: Scalability of Threaded Applications .....</b>	<b>85</b>
<b>Workload-dependent Scaling - An Overview .....</b>	<b>87</b>
Introduction to Sample Application .....	87
<b>1a: Build the Application .....</b>	<b>89</b>
<b>1b: Characterize Multithreaded Implementation with Large Data Set .....</b>	<b>90</b>
<b>1c: Examine Thread Performance with Large Data Set .....</b>	<b>91</b>
<b>1d: Characterize Multithreaded Implementation with Small Data Set .....</b>	<b>92</b>
<b>2: Illustrate Effects of Load Balance in Multithreaded Implementation .....</b>	<b>94</b>
<b>3: Measuring Synchronization Object Overhead.....</b>	<b>95</b>
<b>4: Measuring Front-side Bus Saturation .....</b>	<b>96</b>
<b>5: Identifying False Sharing in Threaded Applications .....</b>	<b>98</b>
<b>Review Questions .....</b>	<b>100</b>
Answers.....	100
<b>Appendices .....</b>	<b>101</b>
<b>A: Setting-up Intel® Thread Checker .....</b>	<b>103</b>
Setting Compile & Link Options .....	103
Creating a New Project and Thread Checker Activity.....	106
<b>B: Setting-up Intel® Thread Profiler .....</b>	<b>109</b>
Setting Compile & Link Options .....	109
Creating a New Project and Thread Checker Activity.....	113
Creating a New Activity by Modifying an Existing Activity .....	116
<b>Additional Resources .....</b>	<b>119</b>







# Lab. 1: Intel® VTune™ Performance Analyzer

<b>Time Required</b>	One hour
<b>Objectives</b>	<p>The sampling profiler in the Intel® VTune™ Performance Analyzer collects system-wide data, has very low overhead, and is easy to use. These labs will familiarize you with the different features and parameters of the profiler and help you interpret its results.</p> <p>At the successful completion of this lab, you will learn to:</p> <ul style="list-style-type: none"><li>• sample-based profiling, VTune analyzer sampling user interface and call graph</li></ul>



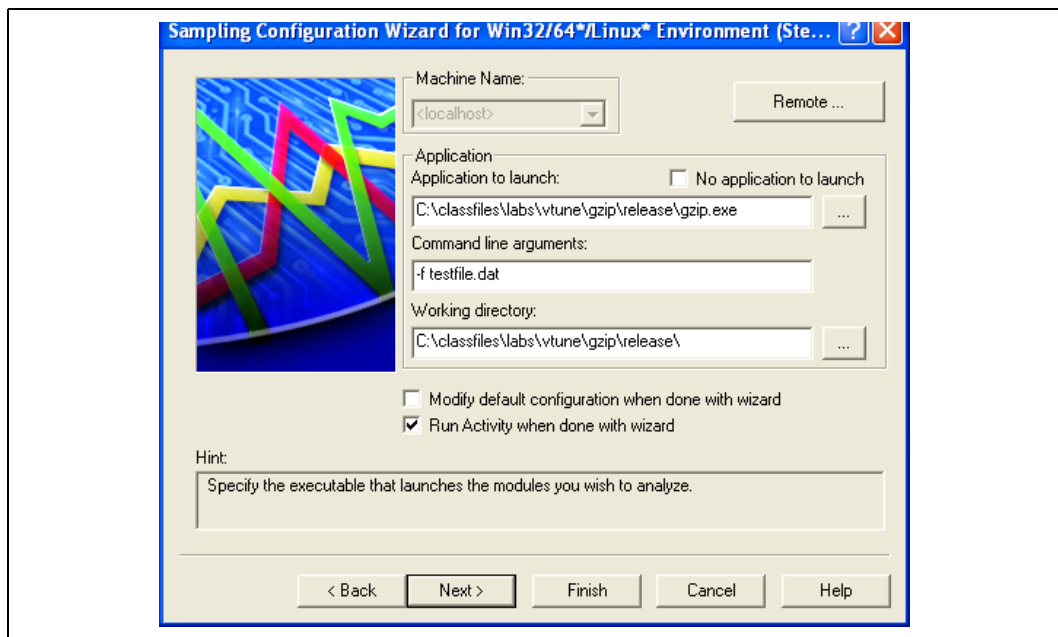
## Activity 1: Finding Hotspots

In this activity, you will collect sampling data for the program gzip.exe and identify the function that consumes the most execution time.

### Collect Sampling Data for the Clockticks Event

1. Make sure that the virus scanner is disabled.
2. Double click the VTune Performance Analyzer icon on the desktop.
3. Click on Category -> Analyzer Projects -> New Project.
4. Click Sampling Wizard.
5. Click OK.
6. Select Window\*/Windows\* CE/Linux Profiling.
7. Uncheck Automatically generate tuning advice.
8. Click Next.
9. In the Application To Launch text box, type the full path of the application, for example:  
`c:\classfiles\VTuneBasics\gzip\release\gzip.exe`
10. In the Command Line Arguments text box, type: `-f testfile.dat`
11. Click Finish. The VTune analyzer will now launch and profile gzip.

Figure 1.1. Sampling Configuration Wizard



## **Review Questions**

- Question 1: What function in gzip.exe takes the most time?
- Question 2: Which function in gzip.exe has the highest CPI?
- Question 3: Which line of source in gzip.exe has the most clocktick samples?
- Question 4: Is gzip.exe multithreaded?

## Activity 2: Sampling Over Time

---

In this activity, you learn how to use the Sampling Over Time view. You analyze the performance of a tbd.

### Create a Sampling Activity

1. Create a new Activity by clicking on Activity->New Activity.
2. Click Sampling Wizard.
3. Click OK.
4. Select Windows\*/Windows\* CE/Linux Profiling.
5. Uncheck Automatically generate tuning advice.
6. Click Next.
7. In the Application To Launch text box, type:

```
c:\classfiles\VTuneBasics\matrix_mt\release\matrix.exe
```

8. Click Finish.
9. After the VTune analyzer finishes collecting data, click the Process button.
10. Click CTRL+A to select all processes.



11. Click the Display Over Time View button, shown at left.

12. You can now see when each of the different processes were running. You can do this for the Process, Thread, and Module views.



13. To zoom in on a particular time region, select the time region by clicking and dragging over it. Then click the Zoom In button, shown at left.



14. To see the regular sampling view for that time region, click the Regular Sampling View button, shown at left. You will now be able to drill down into your code for that time region.

## Activity 3: Using Call Graph

---

In this activity, you will collect call graph data for the program gzip and identify the function that takes the most time and the functions that call it.

### Collect call graph data

1. Create a new Activity by clicking on Activity->New Activity.
2. Double click Call Graph Wizard.
3. Navigate to the classroom directory and locate the gzip.exe application.
4. Click Next.
5. Enter the full pathname for the application, for example:  
c:\classfiles\VTuneBasics\gzip\release\gzip.exe in the Application to launch text box.
6. Type -f testfile.dat in the Command Line Arguments text box.
7. Click Finish.

**Question 1:** What function has the most time spent in it and what functions are calling it?

## **Activity 4: Using the Windows Command Line Interface**

---

In this activity, you learn how to use the Windows\* Command Line Interface to profile an application.

### **Collect Samples based on Clockticks Event using Command Line Interface**

1. Open the command prompt by clicking Start->Programs->Accessories->Command Prompt.
2. Go to the following directory by typing: `c:\classfiles\VTuneBasics\gzip\release\`
3. Create a new Activity that will launch gzip and collect samples based on the Clockticks event by typing: `vtl activity gzip -c sampling -app gzip.exe, "-f testfile.dat"`
4. Run the last Activity created by typing `vtl run`.  
(Alternatively, you can append the word "run" to the end of the command in Step 3 to run the Activity immediately after it is created).

### **View the Profiling Data for gzip**


1. Type: `vtl view -modules` to see the number of samples for each module system-wide.
2. Type: `vtl view -hf -mn gzip.exe` to see the function level breakdown of the samples in gzip.exe.

### **Pack the Data and View it in the GUI**

1. Type `vtl pack gzip`.
2. This creates a file called `gzip.vxp`. This file can be transported from computer to computer and also opened in the VTune analyzer's GUI.
3. Start VTune Performance Analyzer (For example, by double-clicking the desktop icon, if present.)
4. Click Browse for Existing File.
5. Open: `c:\classfiles\VTuneBasics\gzip\release\gzip.vxp`
6. Click OK.
7. Click OK.  
The GUI displays the performance data.







## Lab. 2: Intel® Math Kernel Library

<b>Time Required</b>	Twenty minutes
<b>Objectives</b>	<p>In this lab, you will learn to use Intel® MKL subroutines and Intel® MKL's multithreading capability.</p> <p>At the successful completion of this lab, you will be able to:</p> <ul style="list-style-type: none"><li>• Build images using functions linked from Intel MKL</li><li>• Invoke Intel MKL's multithreading capability</li></ul>



## Activity 1: Matrix Multiply Sample

---

This activity demonstrates performance characteristics of BLAS levels 1, 2 and 3, compared to C source code. In this activity you will inspect, build and run a matrix multiply sample using source code, DDOT, DGEMV and DGEMM.

1. Navigate to the folder MKL\_Overview\DGEMM, and open the file mkl\_lab\_solution.c using any editor of your choice. Go through the code quickly to confirm the 4 implementations of matrix multiply.
2. Examine the Makefiles supplied, and identify the key link steps to enable using MKL. If necessary, edit the file so that all include or library paths are correct. Use these Makefiles to build the demo. Note differences in timings among the different implementations.
3. MKL functions assume a default threads value of 1. Change this value by setting the environment variable OMP\_NUM\_THREADS, for example:

```
set OMP_NUM_THREADS = 2
```

4. Observe the performance at different numbers of threads

**Question 1:** What happens then that the number exceeds the physical processors?

## Activity 2: Monte Carlo Calculation of Pi

---

In this activity, you will modify the Monte Carlo computation of pi to use the random number feature of the Vector Statistical Library (VSL). You will also make use of the multithreading capabilities of VSL.

1. Navigate to the folder MKL\_Overview\MonteCarloPi, and open the file pimonte.c using any editor of your choice. Go through the code quickly to understand how the rand() function is implemented.

**Question 1:** Thought question: Could this loop be threaded?

2. Examine the file pimonte\_VSL.c, and understand the changes necessary to implement a library call to replace rand().

**Question 2:** Why is this not a 1:1 substitution for rand()?

**Question 3:** What is the purpose of, and sensitivity to, blocksize?

**Question 4:** What are the parameters BRNG and VSL\_BRNG\_MCG31?

**Question 5:** Are they the best choices for this computation?

**Question 6:** Could this implementation be threaded?

Examine the Makefiles supplied, and identify the key differences. Use these Makefiles to build versions of this application with rand() and with VSL (recall the syntax "make -f"). Note the impact of the "-up" switch in the compiler report, in the two versions. Note also the difference in results for pi, and for the run times of each image.



## Lab. 3: Programming with Windows\* Threads

<b>Time Required</b>	Twenty-five minutes
<b>Objectives</b>	<p>In this lab session, you will practice multi-threaded programming of applications using Win32* threading API. Sample programs will create threads, assign work to those threads, synchronize access to shared resources, and coordinate thread execution.</p> <p>After successfully completing this lab's activities, you will be able to:</p> <ul style="list-style-type: none"><li>• Find and resolve a common data race involving parameter passing</li><li>• Convert a serial application to a threaded version by encapsulating computations into a function that is executed by threads</li><li>• Find simple data races in code and resolve these threading errors using critical sections as the synchronization mechanism</li></ul>



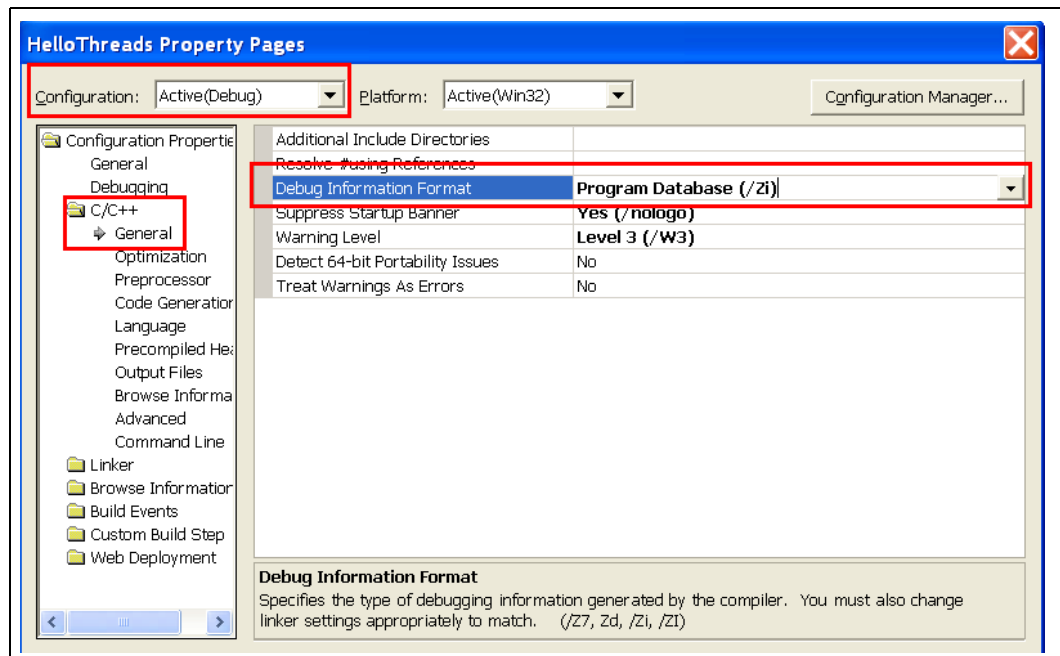
## Activity 1: Starting with *HelloThreads*

<b>Time Required</b>	Ten minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Find and resolve a common data race involving parameter passing</li> </ul>

### Build and Run HelloThreads Program

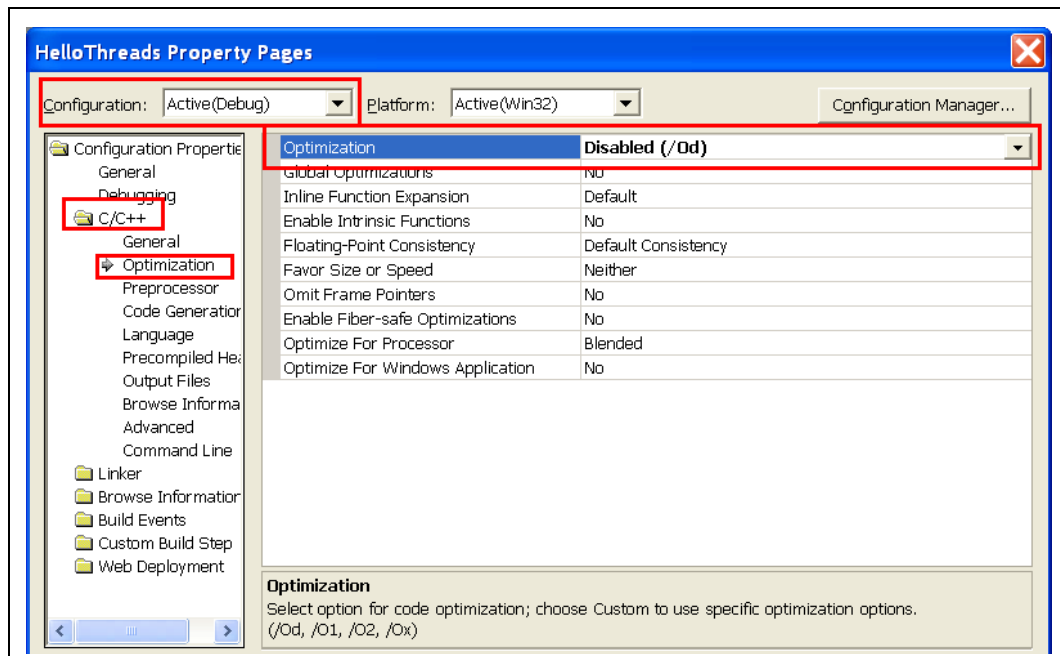
1. Close Microsoft Visual Studio, if it is started.
2. With Windows Explorer\*, open the folder C:\classfiles\Win32 Threads\HelloThreads\.
3. With Microsoft Visual Studio, open the file HelloThreads.sln by double-clicking it.
4. From the Build menu, select Configuration Manager and then select Debug build.
5. From the Project menu, click Properties and then click the C/C++ folder.
6. Make sure that Debug options are selected, as shown in Figure 3.1.

**Figure 3.1. Project Setting – C/C++ Folder – Debug Options**



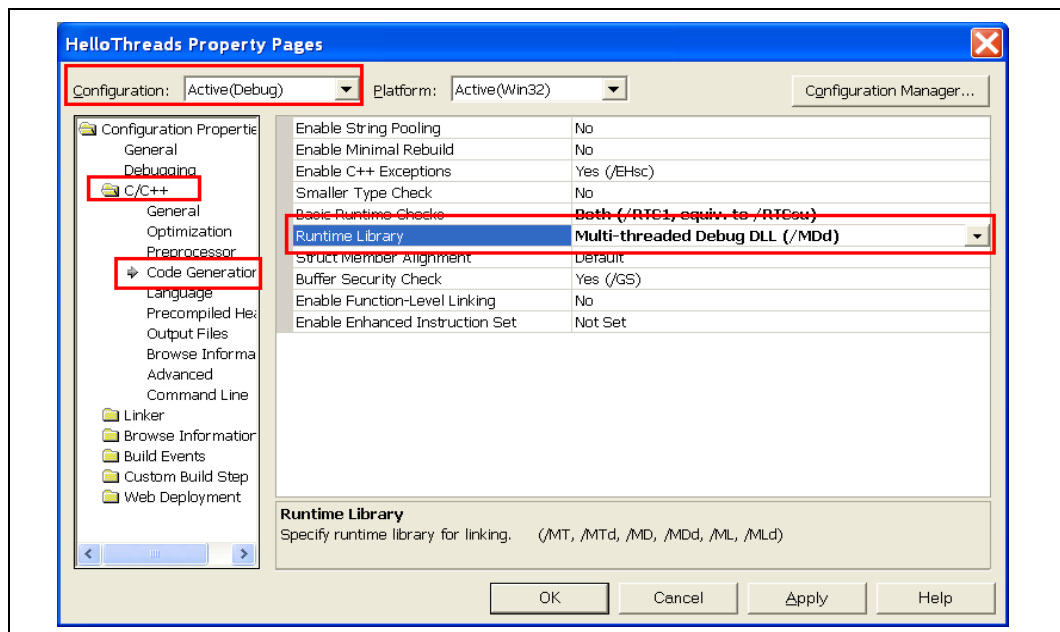
7. Make sure that Optimization is disabled, as shown in Figure 3.2.

**Figure 3.2. Project Settings – C/C++ Folder – Optimization Options**



8. Make sure that thread-safe libraries are selected, as shown in Figure 3.3.

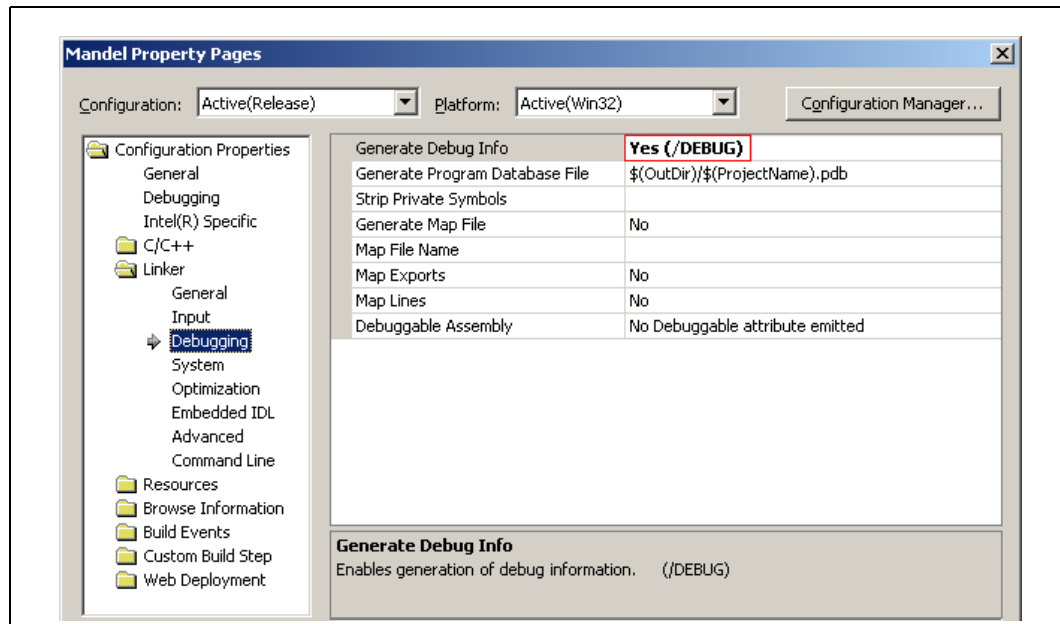
**Figure 3.3. Project Settings – C/C++ Folder – Thread-safe Libraries Options**





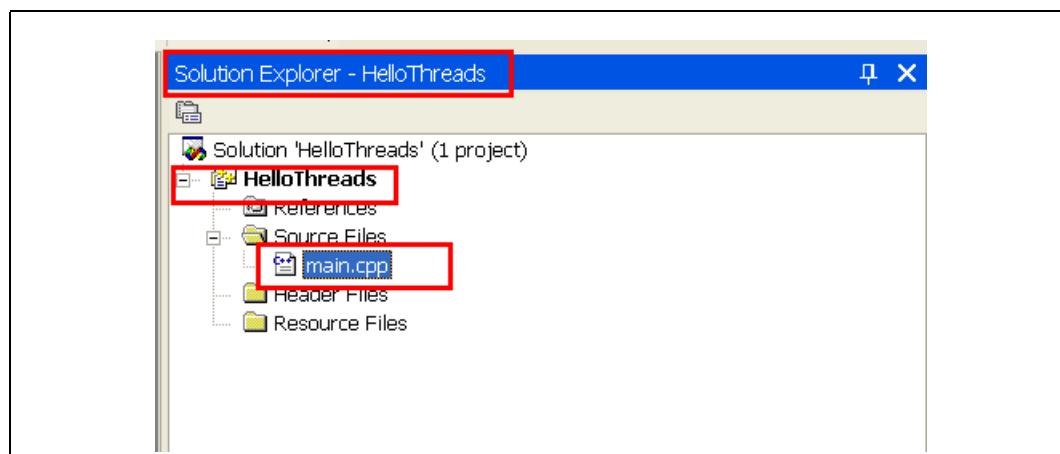
9. Make sure that Debug symbols are preserved during the link phase, as shown in Figure 3.4.

**Figure 3.4. Linker Settings – Debugging Folder**



10. From the Build menu, select Build Solution to build your project.
11. From the Debug menu, select Start Without Debugging to run the program.
12. In Microsoft Visual Studio's Solution Explorer, expand the HelloThreads project and select the file main.cpp to open it, as shown in Figure 3.5.

**Figure 3.5. Solution Explorer – HelloThreads Project**



13. Modify the thread function to report the thread creation sequence (that is, "Hello Thread 0", "Hello Thread 1", "Hello Thread 2", and so on).

*Tip:* Use the CreateThread() loop variable to give each thread a unique number.

14. Build and execute your program.

In what order do the threads execute?  
Do the results look correct?  
Why or why not?

## **Review Questions**

**Question 7:** The execution order of threads is unpredictable.

True                      False

**Question 8:** What build options are required for any threaded software development?

## Activity 2: Approximating $\pi$ with Numerical Integration

<b>Time Required</b>	Fifteen minutes
<b>Objectives</b>	<ul style="list-style-type: none"> <li>• Thread the numerical integration application.</li> <li>• Find any data races in the code and resolve the threading errors using critical sections as the synchronization mechanism.</li> </ul>

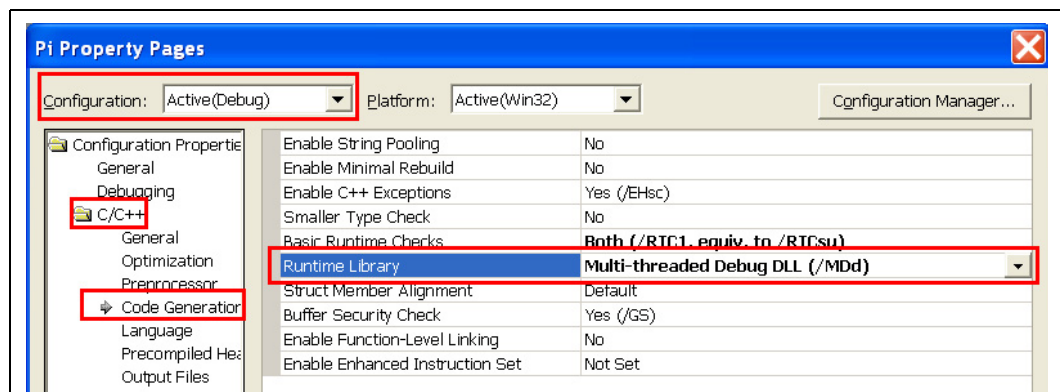
### Build and Run the Serial Program

1. With Windows Explorer, open the folder C:\classfiles\Win32 Threads\Pi\.
2. Start Pi.sln by double-clicking it.
3. From the Build menu, select Set Active Configuration and then select Debug build.
4. From the Project menu, click Properties and click the C/C++ folder. From the Build menu, select Build Solution to build your project.
5. From the Debug menu, select Start Without Debugging to run the program.  
Is the value of the Pi (3.1415...) printed correct?  
Why or why not?

### Correct Errors and Validate Results

1. In Microsoft Visual Studio's Solution Explorer, expand the Pi project and open the file, Pi.cpp.
2. On the C/C++ folder, make sure that thread-safe libraries are selected, as shown in Figure 3.6.

**Figure 3.6. Project Settings – C/C++ Folder – Thread-safe Libraries Options**



3. Thread the serial code to compute Pi using four threads. The bulk of the computation being done is located in the body of the loop. Encapsulate the loop computations into a function and devise a method to ensure that the iterations are divided amongst the threads such that each iteration is computed by only one thread.
4. Use a CRITICAL\_SECTION to protect shared resources accessed by more than one thread. Locate any data races in your program and correct those errors. Some logic changes from the serial version might be required to create code that is both correct and safe.

**Challenge:** Minimize the number of lines of code in the critical section(s).

**Tip:** Think about local variables.

5. When you have changed the source code, from the Build menu, select Build Solution to rebuild and then from the Debug menu, select Start Without Debugging to execute.
6. Keep correcting your source code until you see the correct value of Pi being printed. The correct value of Pi is 3.14159.

**Tip:** A complete solution to the lab is provided in the file PiSolution.cpp, which is located in the following directory: C:\classfiles\Multi-Core\Windows\Win32 Threads\Pi\

## Review Questions

**Question 1:** All threads should use the same CRITICAL\_SECTION.

True                      False

**Question 2:** Threading errors in software can always be corrected by using only synchronization objects.

True                      False

**Question 3:** CRITICAL\_SECTION objects should always be declared as global variables.

True                      False

**Question 4:** What build options are required for *any* threaded software development?

## Activity 3: Using Events

---

<b>Time Required</b>	Fifteen minutes
<b>Objective</b>	Use Windows events to signal when computation threads have completed the assigned work.

The application computes an approximation of the natural logarithm of  $(1 + x)$ ,  $-1 < x \leq 1$ , using the Mercator series. Compute threads are created in a suspended mode. These threads are released to compute the series elements that have been assigned after a "master" thread has been. The master thread waits on a thread count variable that is incremented by each compute thread as it finishes. Once all threads have completed computing the partial sums, the master thread does a final summation and terminates. Results are printed by the process thread after cleaning up all the objects and handles.

### Build and Run Original Threaded Program

1. With Windows Explorer\*, open the folder C:\classfiles\Win32 Threads\N\_ThreadEvent.
2. Double-click on the N\_ThreadEvent.sln icon.
3. Within Microsoft\* Visual Studio, examine the source code until you understand how the threads are created and interact with the others.
4. Build the solution and run the application using the Start Without Debugging command from the Debug menu.

### Modify Original Threaded Program to Use Events

1. Modify the threaded code to use events to signal the master thread rather than the count of a global variable. Some hints for how to do this are given below:
  - a. Create an array of event handles, one event per thread. This can be done like the thread handle array is allocated. Be sure to initialize each event.
  - b. Replace the spin-wait in the master thread with a wait on all the events being signaled from each compute thread.
  - c. Replace the protected increment of the thread count variable with a signal of the event in the array that is indexed with the thread number.  
(Do you still need the critical section?)
2. Build the solution and run the application using the Start Without Debugging command from the Debug menu.

## Activity 4: Using Semaphores

<b>Time Required</b>	Ten minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Identify global data accessed by threads; resolve data races using binary semaphores</li> </ul>

The application opens an input text file. Threads read in lines from the text file and count the total number of words in the line, as well as the number of words with an even number of letters and an odd number of letters. When done, the text file is closed and the final totals are printed.

### Build and Run Serial Program

1. With Windows Explorer, open the folder C:\classfiles\Win32 Threads\SemaphoreLab\.
2. Double-click on the SemaphoreLab.sln icon. Within Microsoft\* Visual Studio, you will find two projects. One is the serial version of the code and the second is the first attempt at threading that code with Win32 Threads.
3. Be sure the Serial project has been selected as the Startup Project. Build the solution and run the serial application using the Start Without Debugging command from the Debug menu.
4. Note the output generated:

Total Words: _____
Total Even Words: _____
Total Odd Words: _____

### Build and Run Threaded Program

1. Set the Threaded project as the Startup Project. Build the solution and run the threaded application using the Start Without Debugging command from the Debug menu. Do you get the same answers as above?.
2. Examine the source code. Identify the global variables that are being accessed by each thread.
3. Rewrite the threaded application to protect the use of these global variables. Wherever mutual exclusion is needed in your solution, use a binary semaphore. Be sure to declare the semaphores at the proper level and initialize them before use.
4. Build and run the threaded code until you are able to achieve the same totals as the serial version of the application.



## Lab. 4: Programming with OpenMP Threads

<b>Time Required</b>	Thirty minutes
<b>Objectives</b>	<p>In this lab session, you will make the Hello World program parallel. Then, you will thread a numerical integration code to compute the value of Pi. Finally, the last lab will have you thread a version of the Pi using a Monte Carlo technique and the Math Kernel Library.</p> <p>After successfully completing this lab's activities, you will be able to:</p> <ul style="list-style-type: none"><li>• Use the most common OpenMP* C statements</li><li>• Compile and run an OpenMP* program</li></ul>





## Activity 1: Starting with *Hello Worlds*

---

In this activity, you make a "Hello, Worlds" program parallel.

### Initial Compile

1. Within the Intel Compiler build environment, navigate to the folder OpenMP\HelloWorlds.
2. Compile serial code using the Intel compiler:

```
icl HelloWorlds.c
```

3. Run the program:

```
HelloWorlds.exe
```

### Add OpenMP Directives

1. Using the editor of your choice (in Windows, the likely candidates will be Notepad or Visual Studio), open the source file and add an OpenMP parallel directive to run the first four lines of main in parallel, but not the last line: `printf("GoodBye World\n");`

```
#pragma omp parallel
{
    ...[Code to run in Parallel goes here]...
}
```

**Question 1:** Do you need to include `omp.h` in the source file? Why or why not?

2. Compile in "Serial Mode" using the Intel compiler:

```
icl HelloWorlds.c
```

3. Notice the pragma warning statements.
4. Fix any syntax errors.
5. Run the program:

```
HelloWorlds.exe
```

### OpenMP Compile

1. Compile in OpenMP mode using the Intel compiler:

```
icl /Qopenmp HelloWorlds.c
```

2. Enable the environment for multiple OpenMP threads:

```
set OMP_NUM_THREADS=2
```

3. Run the program in a multithreaded environment:

```
HelloWorlds.exe
```

4. Run the program multiple times, verifying if the results are the same every time.

## Activity 2: Computing Pi with Numerical Integration

---

In this activity, you will make the Pi program parallel.

### Initial Compile

1. Navigate to the folder OpenMP\Pi.
2. Open the solution file pi.sln (double-click on the file).
3. Build and run this project to ensure that this serial version runs correctly.
4. Record elapsed time: \_\_\_\_\_.

### Add OpenMP Directives

1. Determine the section of code to make parallel and add the OpenMP parallel directive:

```
#pragma omp parallel
{
    ...[Code to run in Parallel goes here]...
}
```

2. Find the loop to make parallel and insert a worksharing pragma:

```
#pragma omp for
for(xxx; yyy; zzz)
{
    //Loop body
}
```

3. Examine all variables and determine which ones need to be specially declared. The following may be handy:

```
#pragma omp parallel private(varname,varname) \
reduction(+:varname,varname) \

shared(varname,varname)
{
    ...[Code to run in Parallel goes here]...
}
```

4. Depending on your implementation you may need the following. For any remaining shared variables add appropriate locks, if you update that variable.

```
#pragma omp critical
{
    ...[Code in Critical section goes here]...
}
```

## OpenMP Compile and Run

1. Update the Project settings to include the compiler flag /Qopenmp
2. Build and run the program.
3. Record the time:\_\_\_\_\_.
4. In a command shell (as in Activity 1), run the program with different settings of OMP\_NUM\_THREADS, and record the time.  
Do the resulting scalings meet your expectations?

## Activity 3: Monte Carlo Pi

---

In this activity, you convert a serial MKL-using version of a Monte Carlo Pi program to a parallel version, using OpenMP directives. The working environment is Visual Studio 2005.

### Initial Compile

1. Navigate to the folder OpenMP\MonteCarloPi.
2. Open the solution file MonteCarloPi.sln (double-click on the file).
3. Build and run this project to ensure that this serial version runs correctly.

*Note:*

You will want to be sure that the MKL include and library paths are set in the project and are correct for the system in use.

4. Record the serial time: \_\_\_\_\_.

### Add OpenMP Directives

1. Determine the section of code to make parallel and add the OpenMP parallel directive:
2. Examine all variables and determine which ones need to be specially declared. There are hints within the source code with regards to some variables and arrays that need to be private to each thread. The following may be helpful:

```
#pragma omp parallel private(varname,varname) \
reduction(+:varname,varname) \

shared(varname,varname)

{

...[Code to run in Parallel goes here]...

}
```

3. Depending on your implementation you may need the following. For any remaining shared variables add appropriate locks, if you update that variable.

```
#pragma omp critical [(name)]

{

...[Code in Critical section goes here]...

}
```

## **OpenMP Compile**

1. Update the Project settings to include the compiler flag /Qopenmp
2. Build and run the program.
3. Record the time:\_\_\_\_\_.
4. In a command shell (as in Activity 1), run the program with different settings of OMP\_NUM\_THREADS, and record the time.  
Do the resulting scalings meet your expectations?

## Review Questions


---

**Question 1:** Name the pragma or directive that would split a loop into multiple threads.

**Question 2:** What website has the very readable OpenMP spec?







## Lab. 5: Correcting Threading Errors with Intel® Thread Checker for Explicit Threads

<b>Time Required</b>	Fifty-five minutes
<b>Objectives</b>	<p>In this lab session, you will learn how to debug Win32* threaded applications and resolve data race conditions. You will learn how to use Intel® Thread Checker, including how to use it with binary instrumentation when Intel® Compilers are not used. Intel Thread Checker is a powerful tool for finding threading errors.</p> <p>After successfully completing this lab's activities, you will be able to:</p> <ul style="list-style-type: none"><li>• Find and resolve data races using simple threading techniques</li><li>• Find actual and potential threading errors, and validate that the errors are fixed using Intel Thread Checker</li><li>• Use Intel Thread Checker to determine if libraries are thread-safe</li><li>• Use the advanced features of Intel Thread Checker, including instrumentation levels</li></ul>



## Activity 1A: Find Prospective Data Races

<b>Time Required</b>	Ten minutes
<b>Objective</b>	<ul style="list-style-type: none"><li>Find and run Intel Thread Checker to find any data races within a simple physics model code</li></ul>

The application computes the potential energy of a system of particles based on the distance, in three dimensions, of each pairwise set of particles. The code is small enough that you may be able to identify the potential data races and storage conflicts by visual inspection. If you identify and make a list of problems, check your list with the list that Thread Checker identifies.

### Build and Run Potential Serial Program

1. With Windows Explorer\*, open the folder C:\classfiles\Thread Checker\potential\_serial
2. With Microsoft Visual Studio, open the file potential\_serial.sln by double-clicking it.
3. From the Build menu, select Configuration Manager and then select Debug build.
4. From the Build menu, select Build Solution (or use Ctrl+Shift+B) and compile the executable.
5. Run the executable (Debug menu -> Start Without Debugging, or Ctrl+F5).

### Build and Run Potential Threaded Program

1. With Windows Explorer\*, open the folder C:\classfiles\Thread Checker\potential\_win
2. With Microsoft Visual Studio, open the file potential\_win.sln by double-clicking it.
3. From the Build menu, select Configuration Manager and then select Debug build.
4. Ensure that the following compile and link flags are set correctly (See [Appendix A: "Setting-up Intel® Thread Checker" on page 103](#) for help):
  - a. Debug options are set (/Zi)
  - b. Debug symbols are preserved during linking (/DEBUG)
  - c. Optimization is disabled (/Od)
  - d. Thread safe system libraries are used (/MDd)
  - e. The binary is re-locatable (/fixed:no)
5. From the Build menu, select Build Solution to compile the executable.
6. From the Debug menu, select Start Without Debugging to run the program.

**Note:** The number of particles and iterations has been reduced from the serial version of the code in order to facilitate Thread Checker runs.

7. With Windows Explorer\*, start the Intel VTune Performance Analyzer.
8. Start a New Project by clicking on the icon.
9. From the Category Threading Wizards, select the Intel® Thread Checker Wizard.
10. Set up the potential\_win.exe application to be run by using the browse “...” button and then click Finish to start Thread Checker.
11. Scan through some of the diagnostics displayed. Check the source code lines from some of the diagnostics.  
Can you see why there is a conflict on those lines of code?

## Activity 1B: Resolve Data Races

---

Time Required	Ten minutes
Objective	<ul style="list-style-type: none"><li>Resolve data races found in previous lab using simple threading techniques</li></ul>

### Resolve the Problems

1. Fix the problems identified by Thread Checker in the earlier lab. Which variables can be left to be shared between threads? Which variables can be made local to each thread? Which variables must be protected with some form of synchronization?
2. Be sure to re-run the corrected application through Thread Checker until you have no more diagnostics being generated from the code.
3. Once you have eliminated all the threading problems, re-set the number of particles and iterations that were used in the serial version. Re-build and run the code.

Do the answers from the threaded code match up with the output from the serial application?

## Activity 2: Identifying Deadlock

---

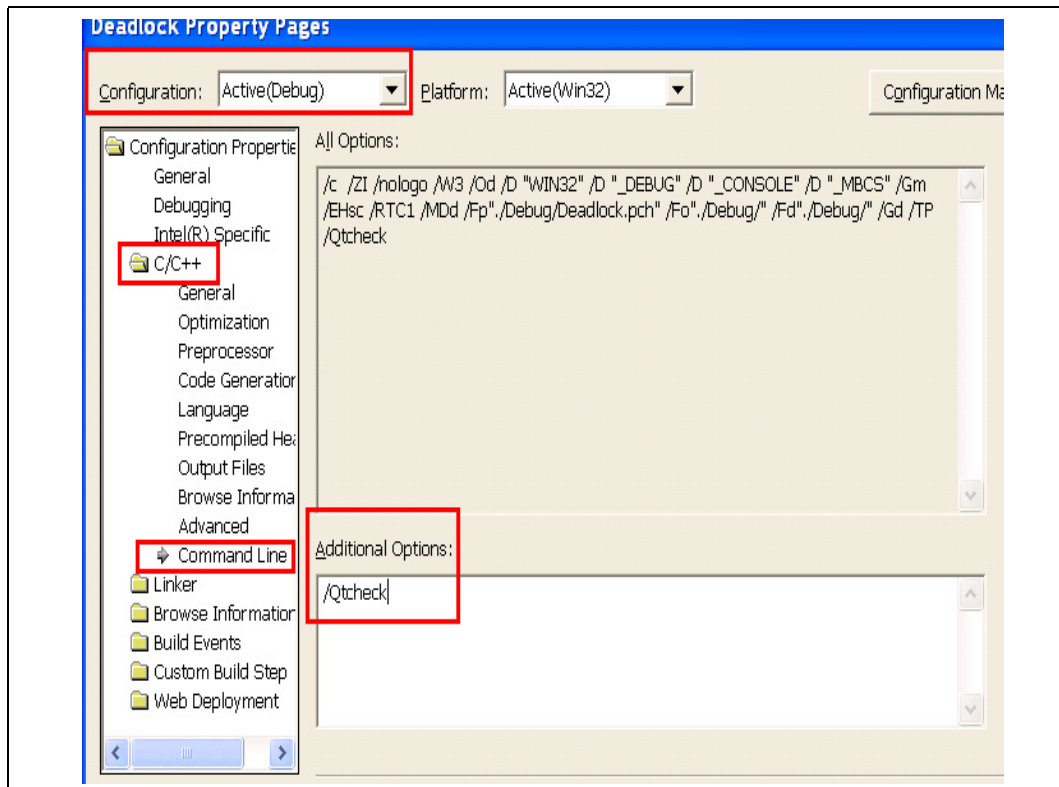
<b>Time Required</b>	Fifteen minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>• Build software to use source code instrumentation for Intel Thread Checker</li> <li>• Find actual and potential threading deadlock errors, and validate that the errors are fixed using Intel Thread Checker</li> </ul>

### Build and Run the Program

1. With Windows Explorer, open the folder C:\classfiles\Thread Checker\Deadlock\.
2. Open Deadlock.sln by double-clicking it.
3. From the Build menu, select Configuration Manager and then select the Debug build.
4. Using the Debug configuration project, ensure that the following compile and link flags are set properly (See [Appendix A: "Setting-up Intel® Thread Checker" on page 103](#) for help).
  - a. Debug options are set (/Zi)
  - b. Debug symbols are preserved during linking (/DEBUG)
  - c. Optimization is disabled (/Od)
  - d. Thread safe system libraries are used (/MDd)
  - e. The binary is re-locatable (/fixed:no)

5. On the C/C++ menu, select the Command Line folder and add /Qtcheck by editing the Additional Options, as shown in Figure 5.1.

**Figure 5.1. Project Settings – C/C++ Folder – Compiler Options**



6. From the Build menu, select Build Solution to build your project.

Is this application using:

Source instrumentation? \_\_\_\_\_

Binary instrumentation? \_\_\_\_\_

How can you tell which instrumentation type is used?

7. From the Debug menu, select Start Without Debugging to run the program.

Do the results look correct?

Why or why not?

## Run Application within Intel® Thread Checker

1. Start VTune Performance Analyzer and select **New Project**. (See [Appendix A: "Setting-up Intel® Thread Checker" on page 103](#) for help.)
2. From the **Category Threading Wizards**, select the **Intel® Thread Checker Wizard**.
3. Set up the application to be run by using the browse "... " button and then click **Finish** to start Thread Checker.  
You should see either one of the two displays shown in [Figure 5.2](#) and [Figure 5.3](#), because deadlock does not occur during every run.

[Figure 5.2](#) shows the diagnostics list when the deadlock does not occur during the run. Intel Thread Checker still catches it as a potential deadlock and it is listed with YELLOW bullets (circled in red), which indicates CAUTION.

**Figure 5.2. Diagnostic List When Deadlock Does Not Occur During a Run**








Context[Best] ▲	ID	Short Description	Severity	Description	Count	Filtered
"Deadlock.cpp":	1	Read -> Write data-race		Memory write at "Deadlock.cpp":43 conflicts with a prior memory read at "Deadlock.cpp":67 (anti dependence)	1	False
"Deadlock.cpp":41	2	Read -> Write data-race		Memory write at "Deadlock.cpp":45 conflicts with a prior memory read at "Deadlock.cpp":67 (anti dependence)	1	False
Whole Program 1	3	Thread termination		Thread Info at "Deadlock.cpp":64 - includes stack allocation of 1048576 and use of 4096 bytes	1	False
Whole Program 2	4	A sync object was acquired in the wrong order		A synchronization object "Deadlock.cpp":44 was acquired in the wrong order at "Deadlock.cpp":42	1	False
Whole Program 2	5	A sync object was acquired in the wrong order		A synchronization object "Deadlock.cpp":31 was acquired in the wrong order at "Deadlock.cpp":29	1	False
Whole Program 3	6	Thread termination		Thread Info at "Deadlock.cpp":63 - includes stack allocation of 1048576 and use of 4096 bytes	1	False
Whole Program 4	7	Thread termination		Thread Info at "Deadlock.cpp":54 - includes stack allocation of 1048576 and use of 4096 bytes	1	False



Figure 5.3 shows the error list when the deadlock occurs during a run. When this occurs, the user **MUST** kill the application by closing the DOS window opened by the application.

**Figure 5.3. Diagnostic List When Deadlock Occurs During a Run**

Context[Best] >	ID	Short Description	Severity	Description	Count	Filtered
"Deadlock.cpp":27	1	Read -> Write data-race		Memory write at "Deadlock.cpp":29 conflicts with a prior memory read at "Deadlock.cpp":67 (anti dependence)	1	False
"Deadlock.cpp":34	2	Read -> Write data-race		Memory write at "Deadlock.cpp":42 conflicts with a prior memory read at "Deadlock.cpp":29 (anti dependence)	1	False
"Deadlock.cpp":34	3	Write -> Read data-race		Memory read at "Deadlock.cpp":42 conflicts with a prior memory write at "Deadlock.cpp":29 (flow dependence)	1	False
"Deadlock.cpp":34	4	Write -> Write data-race		Memory write at "Deadlock.cpp":42 conflicts with a prior memory write at "Deadlock.cpp":29 (output dependence)	1	False
"Deadlock.cpp":	9	The application was forcefully...		The application was forcefully terminated	1	False
Whole Program 1	5	A Thread is deadlocked		A Thread at "Deadlock.cpp":68 is deadlocked trying to acquire a resource owned by a thread at "Deadlock.cpp":64	1	False
Whole Program 2	6	A Thread is deadlocked		A Thread at "Deadlock.cpp":68 is deadlocked trying to acquire a resource owned by a thread at "Deadlock.cpp":63	1	False
Whole Program 3	7	A Thread is deadlocked		A Thread at "Deadlock.cpp":43 is deadlocked trying to acquire a resource owned by a thread at "Deadlock.cpp":28	1	False
Whole Program 4	8	A Thread is deadlocked		A Thread at "Deadlock.cpp":30 is deadlocked trying to acquire a resource owned by a thread at "Deadlock.cpp":41	1	False
Whole Program 5	10	Thread termination		Thread Info at "Deadlock.cpp":47 - includes stack allocation of 1048576 and use of 4096 bytes	1	False
Whole Program 6	11	Thread termination		Thread Info at "Deadlock.cpp":63 - includes stack allocation of 1048576 and use of 4096 bytes	1	False
Whole Program 7	12	Thread termination		Thread Info at "Deadlock.cpp":64 - includes stack allocation of 1048576 and use of 4096 bytes	1	False

- Double-click on one or more of the deadlock diagnostics (error or warning) in the **Diagnostics** window. Explore what lines Thread Checker points out as being involved in the deadlock diagnostics.

## Correct Errors and Validate Results

- Return to Microsoft Visual Studio to correct the program.
- From Microsoft Visual Studio, edit the program and re-build it.

*Tip:*

Notice the order that the CRITICAL\_SECTION objects cs0 and cs1 are used in each thread function work0() and work1().

- When you have changed the source code, from select Build Solution the Build menu to rebuild.
- In the VTune analyzer environment, click the Activity menu and select Run to run Intel Thread Checker again to validate that you have corrected all threading errors.

You can re-use the same VTune analyzer environment project; in fact, have the VTune analyzer environment started and running while you change the source code in Microsoft Visual Studio.

Once you have no threading errors, the display will contain only informational messages.

*Tip:* Minor logic changes will be required to correct all threading errors.

A complete solution to the lab is provided in the file, DeadlockSolution.cpp.

## **Review Questions**

1. What build options are required for *any* threaded software development?
2. What build options are required for binary instrumentation?
3. What build options are required for source instrumentation?
4. Under what category is the Intel Thread Checker found in the VTune analyzer environment?
5. Using a larger data set (workload) causes Intel Thread Checker to find more information (errors, warnings, and so on). True False
6. Threading errors in software can always be corrected by using only synchronization objects. True False
7. If a deadlock is present in an application, it will always occur at run time. True False

## Activity 3: Testing Libraries for Thread Safety

Time Required	Twenty minutes
Objective	<ul style="list-style-type: none"><li>Use the advanced features of Intel Thread Checker, including instrumentation levels</li></ul>

### Setting-up and Compilation for Thread Safety Testing

1. With Windows Explorer, open the folder C:\classfiles\Thread Checker\Thread Safe Libraries\.
2. Double click the Thread Safe Libraries.sln Microsoft Visual Studio solution file to open the projects for this lab. You will find two projects in this solution:
  - a. Library – this project will build a DLL containing some functions to be tested for thread safety. This DLL can be thought of as either a user library or a third-party library.
  - b. Library Tester – this project is used to call the library functions and will be modified to test the library routines for thread safety.
3. Using the Debug configuration, build and run the Library Tester project to better understand the library routines. (This will also build the Library DLL.)

**Note:**

Since we will be using OpenMP to generate threaded calls to the library routines involved and Intel Thread Checker to determine if there are any data conflicts or other problems, you must first convert the Library Tester project to use the Intel compiler.

4. Right-click on the Library Tester Project and select "Convert to use Intel[R] C++ Project System." Click "Yes" in the Confirmation dialog box.
5. Add some OpenMP parallel sections code to run each of the 6 pairwise combinations of the three library routines on threads.  
For example, to test the thread safety of a library routine, say foo(), with itself, use the following code:

```
#pragma omp parallel sections
{
    #pragma omp section
        foo(x);
    #pragma omp section
        foo(y);
}
```

6. Before building the threaded library testing application, be sure to set the compilation to use OpenMP. Within the C/C++ configuration properties folder of the Language category, you will see a sub-pane labeled "Intel Specific". Under this, change the "Process OpenMP Directives" item to "Generate Parallel Code (/Qopenmp)."
7. Rebuild the application.
8. After launching VTune, create a new project and define a new Thread Checker activity for the library\_tester.exe application. Run the Thread Checker activity.

Were any diagnostics found?

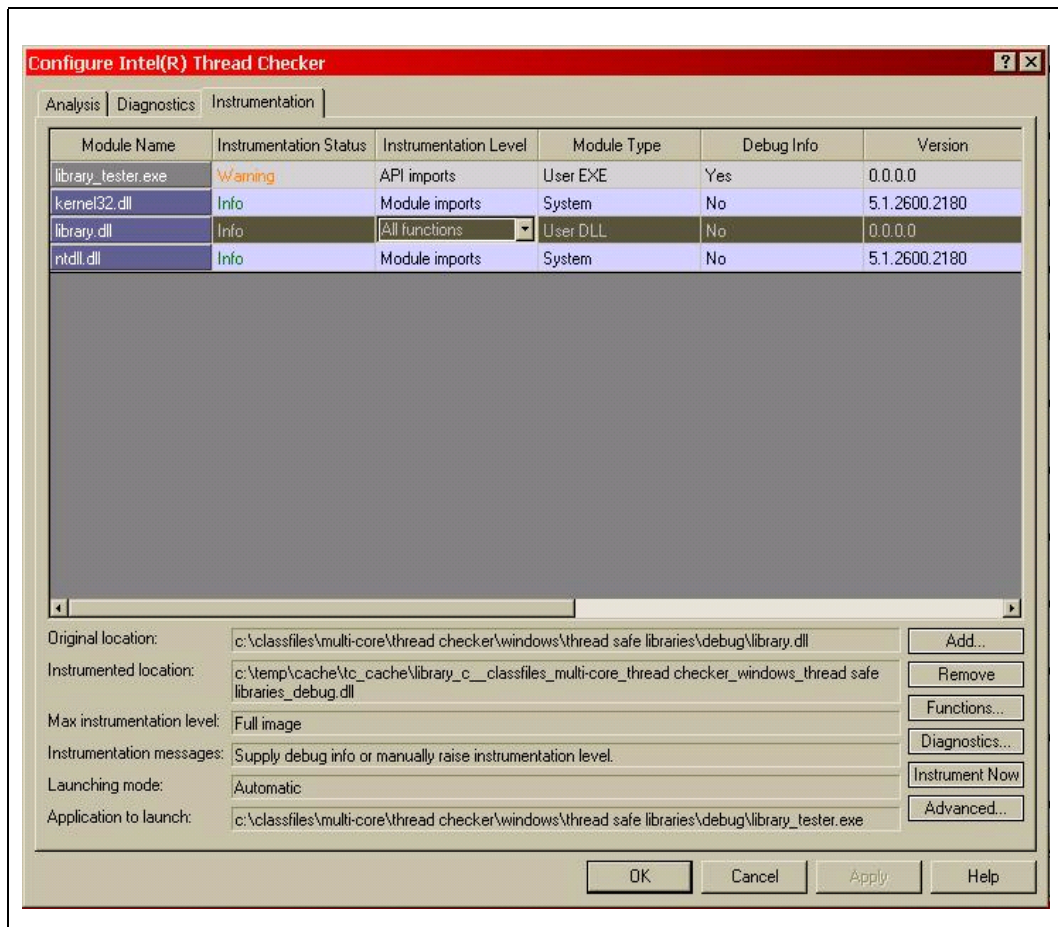
## Modifying Binary Instrumentation Levels

Even though the library DLL was compiled with debug symbols and information, Thread Checker may not find any diagnostics - even if there are data conflict errors within the library. This is due to the level of binary instrumentation used by Thread Checker.

The default binary instrumentation for user DLLs is "All Functions," which will instrument each instruction of those portions that have debug information. However, Thread Checker is unable to locate the debug information data base file and lowers the instrumentation level to "API Imports". This reduced level will not instrument user code, but will instrument selected system API functions. Thus, to ensure that the library functions are thread safe, we must manually raise the instrumentation level of the DLL and re-run the analysis.

1. In the Thread Checker "Tuning Browser" pane, right click on the activity and chose the "Modify Collectors" option.
2. In the dialog box that pops up select the "Modify the selected Activity" radio button and click OK.
3. Click on the Instrumentation tab in the "Configure Intel Thread Checker" dialog box(Figure 5.4).
4. Left-click on the "Instrumentation Level" entry for the "library.dll" row. Select "All Functions" from the pull-down menu. In the lower right corner, click on the "Instrument Now" button to perform the binary instrumentation of this DLL at the chosen level.

Figure 5.4. Setting binary instrumentation levels for DLLs and other modules



5. Click OK.
6. Re-run the Thread Checker analysis.  
Were any diagnostics found this time?

## Review Questions

- Question 1:** Which combinations of functions are not thread safe?
- Question 2:** If this were a third-party library, rather than a set of function to which you have source code access, can you determine which functions were not thread safe in order to report this problem to the library developers?





## Lab. 6: Tuning Threaded Code with Intel® Thread Profiler for Explicit Threads

<b>Time Required</b>	Forty-five minutes
<b>Objectives</b>	<p>In this lab session, you will use Intel® Thread Profiler to detect performance issues in applications threaded with the Win32 threading API.</p> <p>After successfully completing this lab's activities, you will be able to:</p> <ul style="list-style-type: none"><li>• Create new activities under the Thread Profiler</li><li>• Use the Thread Profiler to determine system utilization</li><li>• Navigate through the Thread Profiler features to understand program thread activity</li></ul>





## Activity 1A: Getting Started with Thread Profiler

---

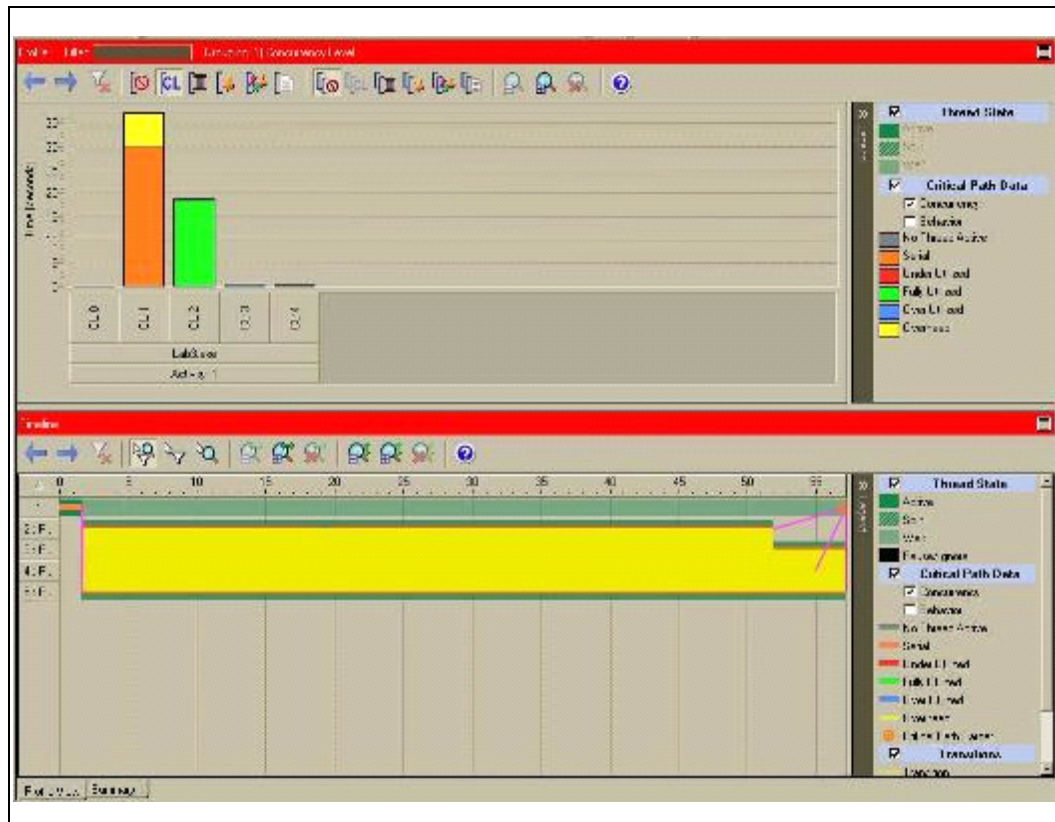
Time Required	Ten minutes
Objective	<ul style="list-style-type: none"><li>Start up Intel Thread Profiler and examine the different analysis views offered by it</li></ul>

The application computes the potential energy of a system of particles based on the distance, in three dimensions, of each pairwise set of particles. This lab is designed to give you an introduction to running an application through the Thread Profiler and seeing what views are available within the tool. You will also see what those views reveal about the execution of threads within an application.

1. Find the Microsoft Visual Studio solution file C:\classfiles\ThreadProfiler\potential lab 1\Potential Lab 1.sln.
2. Double-click the icon to open Microsoft Visual Studio and build the application. (See [Appendix B: "Setting-up Intel® Thread Profiler" on page 109](#) for compilation and linker settings needed for the Thread Profiler analysis.)
3. Launch VTune Performance Analyzer and create a new Thread Profiler activity that will run the executable generated from the previous step.
4. Run the application from within the Thread Profiler.

5. On completion of the run, you should see the split Profile View and Timeline View as shown in Figure 6.1.

**Figure 6.1. Thread Profiler Default View Screen**



6. Click on the expansion button in the upper right corner of the Profile View. You will see the Concurrency Level button in the toolbar has been automatically selected.
7. Click on the other grouping buttons ("Threads View" and "Object View"). Click on the bars and roll over portions of the windows to see what tooltips appear and what data is contained within them.
8. Click on the expansion button of the Timeline Tab to see the "Timeline View" of the performance data. Explore the window displayed and click on parts of the display to see what information might be available from this view.

## Review Questions

- Question 1:** From what you've seen, how would you characterize the performance of this application?
- Question 2:** Are there any obvious performance issues?

## Activity 1B: Analyzing an Application

<b>Time Required</b>	Ten minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Examine the different analysis views offered by Intel Thread Profiler and determine if a performance issue is evident from the information presented</li> </ul>



- Return to the Profile View and bring up the “Concurrency Level” view.

Approximately, what percentage of time was spent in serial (one thread only) and under-subscribed execution on the platform? \_\_\_\_\_

What percentage of time (approximately) was spent in full parallel execution?  
\_\_\_\_\_

What is this view telling you? \_\_\_\_\_



- Select the “Threads View” by clicking on that button.

This view shows all the threads in the applications and how each of those were active in this critical path.

How many total threads were used during the execution of this application?

Can you tell if there is some performance problem inherent in the application?

Yes      No

If so, what problem is it? If not, why not?



- Select “Objects View” by clicking on that button.

Are there any synchronization objects that account for a significant portion of the critical path time spent in serial or under-subscribed impact time?

If so, which one(s)?

- Select Expand the “Timeline” to bring up the Timeline View.

What is the most striking feature that you notice from this data?

Can you tell if there is some performance problem inherent in the application from this view?

Yes      No

If so, what would you suggest be done to correct this problem?

## Activity 2: Finding Load Balance Issues

<b>Time Required</b>	Fifteen minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Use Intel Thread Profiler to find a load imbalance performance problem within a threaded application</li> </ul>

The application used for this lab is a version of the potential energy physics simulation that employs a pool of threads, rather than creating and terminating a new set of threads for each time-step. Threads are controlled by a pair of events that signal threads to start and also signal the main thread when threaded execution has completed. Even though the loop iterations over the particles has been divided equally between threads, there is still a load balance issue with this code.

### Build and Run Potential Threaded Program

1. With Windows Explorer\*, open the folder  
C:\classfiles\ThreadProfiler\Potential Lab 2\.
2. Double-click the Potential Lab 2.sln icon to start Microsoft Visual Studio.
3. Notice in the source code, the definition and initialization of the two events (bSignal and eSignal). Also, notice the use of the *done* variable within the `tPoolComputePot` function. This controls the termination of the threads when the simulation has completed.
4. Select the Release build, ensure the debug and linker options are set as needed, and build the application.
5. Launch VTune Performance Analyzer and create a new Thread Profiler activity that will run the executable generated from the previous step.
6. Run the application from within the Thread Profiler.

### Evaluate Thread Profiler Results to Diagnose Problem

The Concurrency Level Profile View confirms that the majority of execution time is spent with only one thread active. If you switch on the Behavior categories for the Critical Path Data, a large percentage of time is spent in Impact time. This data means that a single thread is running and actively keeping other threads from executing, typically through some synchronization object.

1. Bring up the Object View.

Can you tell if there is one object that has been involved with the Serial Impact Time? If so, which one?

Is this object protecting access to data?

If so, can the application and data access patterns be modified to reduce the amount of time spent by threads "holding" this synchronization object?

**Note:** Since the synchronization object (eSignal) is used to signal when the threaded portion of the computations are completed, and not to protect data, the main thread is the only one "impacted." This impact is necessary to ensure the correct execution of the application. Thus, modifications to the data will not affect this situation.

2. If the problem is not with data-protecting synchronization, the cause may be within the threads themselves. Bring up the Threads profile view.

Can you tell if one or more threads were involved in the "Serial impact time" from the critical path? If so, which ones?

3. Compare the Lifetime and Active time of each of the `tpoolComputePot` (worker) threads.

Thread 2 Lifetime: _____	Thread 2 Active time: _____
Thread 3 Lifetime: _____	Thread 3 Active time: _____

What does this tell you about the application's threaded execution?

Should we be concerned with the Serial cruise time within Thread 1?

## Fixing the Performance Issue

**Note:** The difference in active times of the worker threads indicates that there is a load imbalance between the amounts of computation assigned to these threads. You will need to reconfigure how work is assigned to each thread, in order to achieve a more balanced amount of work between the worker threads.

1. Bring up the source code within the Microsoft Visual Studio window.
2. Notice that the first loop within the main routine statically divides up the particle iterations based on the number of worker threads that will be created within the thread pool.

In the `computePot` routine, each thread uses the stored boundaries indexed by the thread's assigned identification number (`tid`), to fix the start and end range of particles to be used. However, the inner loop within this routine uses the outer index within the exit condition. Thus, the larger the particle number used in the outer loop, more iterations of the inner loop will be executed. This is done so that each pair of particles contributes only once to the potential energy calculation.

There are two obvious ways to fix this load imbalance:

- I. Because the computation for each pair of particles considered will be equivalent, modify the code to:
  - i. find the number of such computations that will be done  $((N*2)/2 - N)$
  - ii. divide this by the number of threads being used
  - iii. compute groupings of particles that will yield the closest set of computations calculated in the previous step, and
  - iv. statically set the bounds array entries to create these groupings to be assigned to threads.
- II. Use a more dynamic assignment of particles to threads. For example, rather than assigning consecutive groups of particles, as in the original version, have each thread, starting with the particle indexed by the thread id `tid`, compute all particles

**Lab. 6: Tuning Threaded Code with Intel® Thread Profiler for Explicit Threads**

whose particle number differ by the number of threads. For example, when using two threads, one thread handles the even-numbered particles while the other thread handles the odd-numbered particles.

The first scheme involves considerable amount of code modifications, but will still be scalable for different numbers of threads and/or number of particles, and still achieve a good load balance. The second scheme will achieve similar results and scalability, but requires much less code modifications.

3. Modify the physics simulation code to achieve a better load balance between the worker threads. You can use one of the solutions outlined above, or use one of your own design.
4. Re-run your modified code through the Thread Profiler to see if you have achieved the results you desired.

*Note:* After making such changes, you would normally run the modified application through the Thread Checker to ensure that no new threading errors had been introduced. If you have achieved sufficient load balance from the modified code and have the time, you should test the application with the Thread Checker.

## Activity 3: Finding Synchronization Contention Issues

<b>Time Required</b>	Twenty minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Use Intel Thread Profiler in order to find a synchronization contention performance issue within a threaded application. Demonstrate some of the filtering capabilities of the Thread Profiler.</li> </ul>

The application computes an approximation to the value of the constant Pi using the "midpoint (rectangle) rule" of numerical integration. The worker threads that are created compute the area of rectangles using the function value at selected points as the height of a rectangle. The results of each height computation are stored into a global sum. Access to this global variable is (correctly) protected by a CRITICAL\_SECTION object. Even though the number of rectangle-area computations has been divided equally between threads, there is still a performance issue with this code.

### Build & Run Threaded Numerical Integration Program

1. With Windows Explorer\*, open the folder  
C:\classfiles\ThreadProfiler\Numerical Integration\.
2. Double-click the TP\_Lab3.sln icon to start Microsoft Visual Studio
3. Select the Release build, ensure the debug and linker options are set as needed, and build the application.
4. Launch VTune Performance Analyzer and create a new Thread Profiler activity that will run the executable generated from the previous step.
5. Run the application from within the Thread Profiler.

### Evaluate Results to Diagnose Problem

On completion of the run, you should examine the Concurrency Level view. There is a large portion of the execution along the critical path that was spent in "Serial impact time." If there is a way to eliminate or reduce this, the application will run faster. Also, there is quite a bit of time being spent in threading overhead.



1. Double click on the critical path bar to go to the profile view. If the Concurrency Level view does not open as the default, click on the "Concurrency Level" button in the toolbar.  
What is this view telling you? \_\_\_\_\_



2. Select the "Threads View" by clicking on that button. This view shows all the threads in the applications and how those were active in this critical path. Is there a load imbalance between the worker threads?



3. Select "Object View" by clicking on that button. This view will show that the application is impacted by one Critical Section object where all of the accumulated

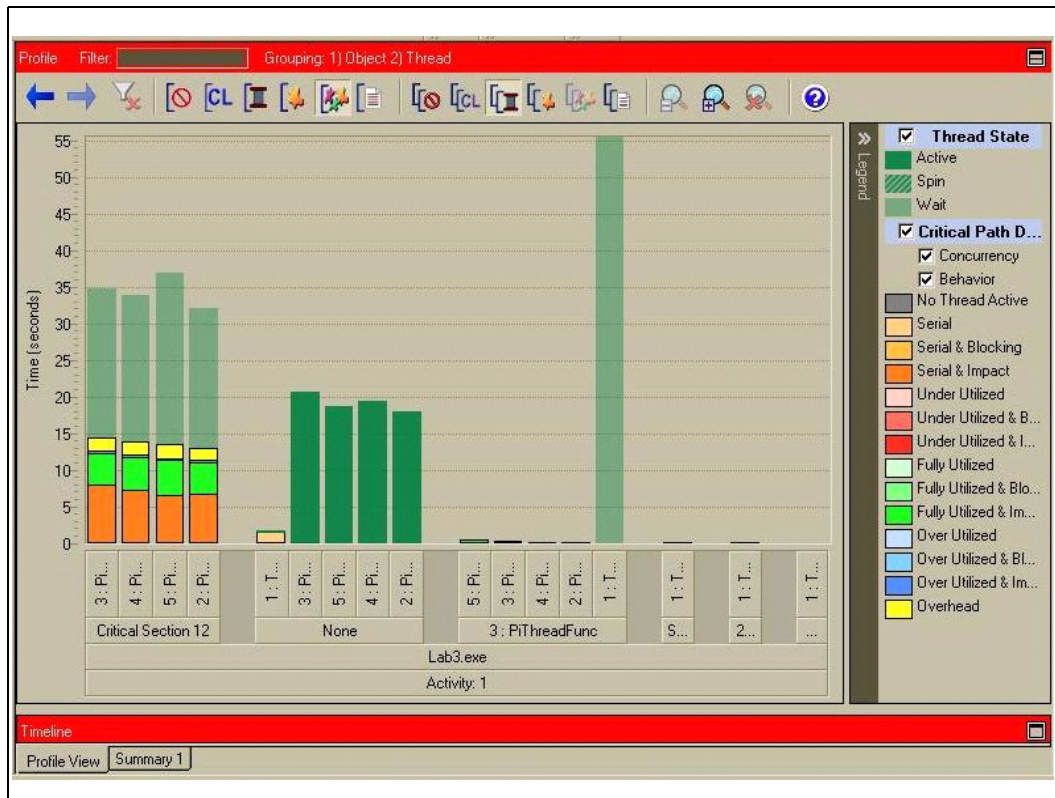


impact is present.



4. Select a second level grouping in the Object View by clicking on that button to choose Threads as the secondary grouping. This will bring up a view similar to the one shown in Figure 6.2.

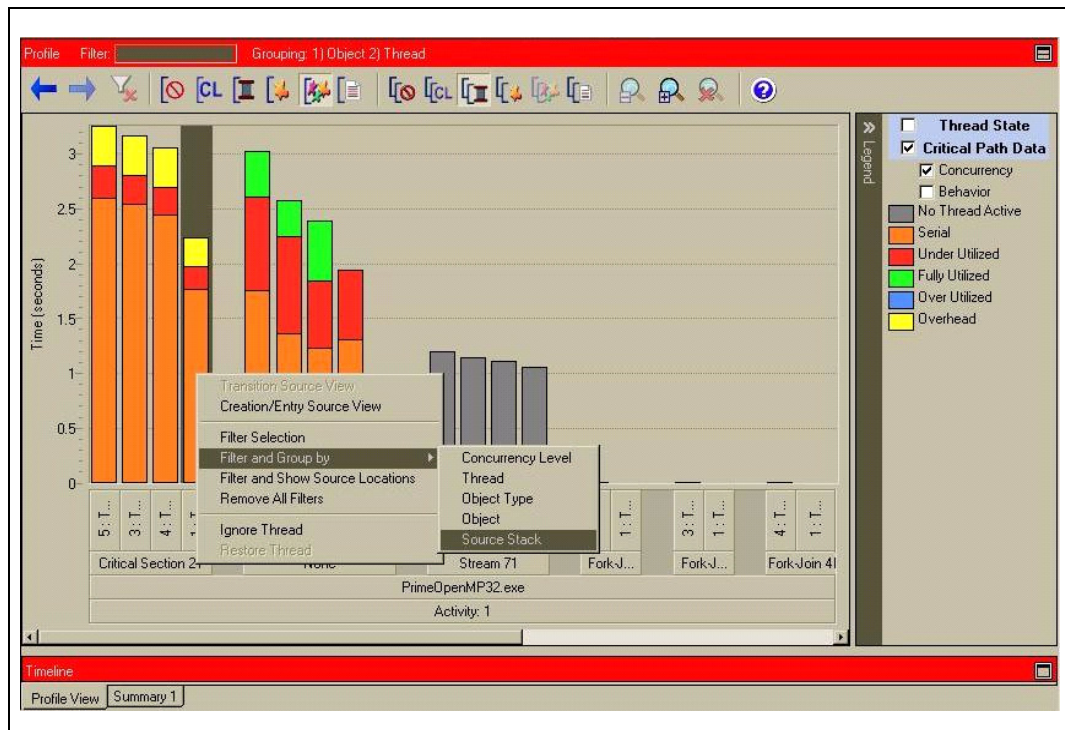
**Figure 6.2. Second Level Grouping under Objects View**



Both worker threads are impacted by the same Critical Section Object. The next few steps show how to get from the impacting object to the source line.

5. Select one of the Threads bar impacted by the critical section object to see the pop-up menu shown in Figure 6.3.

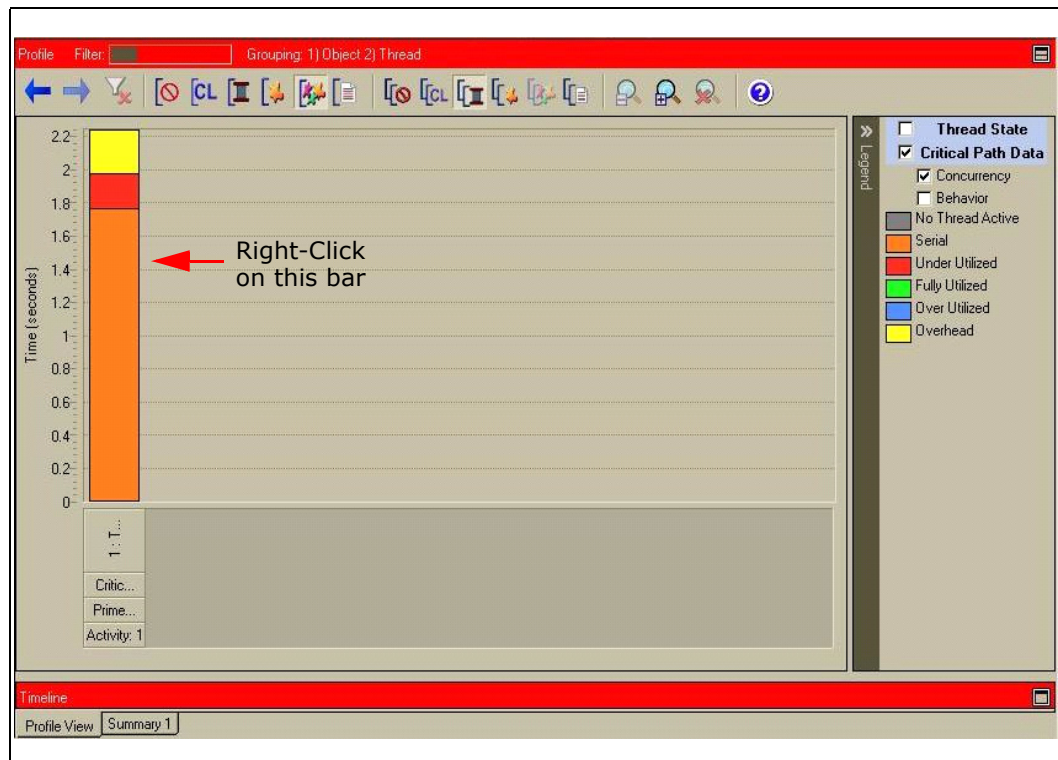
Figure 6.3. Filtering by Object and Grouping by Source Code Locations



6. Select "Filter and Group by" drop down menu and in the drop-down, select Source Stack and in the drop down, and then select All.

This should bring up a screen similar to the one shown in [Figure 6.4](#).

Figure 6.4. Filtered Profile View

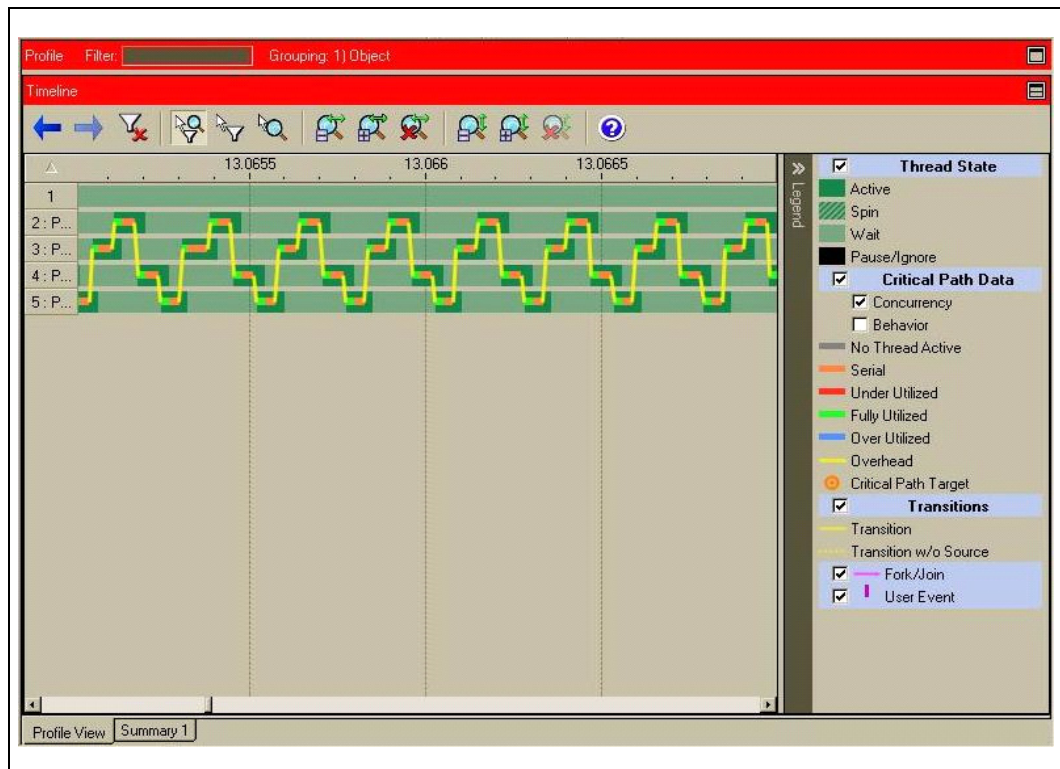


7. Right click on the bar graph and select "Transition Source View" in the selection. This above operation will take you to the source location where the object is used and shows the transition information.

## Using the Timeline View for Analysis

1. Return to the Profile view. Remove the filtering by clicking on the <funnelX> icon. (Though it is not necessary, you may also want to go back to a simple display by clicking the <Obar> icon.)
2. Click on the Timeline expansion button to bring up the Timeline view. If you did not have "Transitions" set when the application was run, modify the collectors (See [Section "Creating a New Project and Thread Checker Activity" on page 113](#)). Re-run the application in the Thread Profiler; click on the Timeline tab after results are displayed.
3. Zoom (click and drag) into a portion near the middle of the timeline that has multiple threads running. Continue zooming in until you have a view similar to the one shown in [Figure 6.5](#).

Figure 6.5. Timeline View (Zoomed)



What is this view telling you? \_\_\_\_\_

4. Hold the pointer over one of the red impact bars. The tooltip box details which synchronization object was involved and identifies the threads involved.
5. Hold the pointer on one of the transition lines. The tooltip popup details the synchronization object, the threads involved and the source lines with the transition of the critical path from one thread to another.
6. Right click on a transition line. Choose "Transition Source View" from the menu popup to open a Transition Source window.  
This will show you the source code location where the critical path transitioned from one thread to another.

## Fixing the Performance Issue

The repeated access and contention on the CRITICAL\_SECTION object causes a large part of the run to be executed in serial. If more worker threads were used, the problem would be worse, as more threads would sit idle waiting to acquire the synchronization object.

1. Modify the numerical integration code to achieve better performance.  
To fix the contention problem, declare a variable in the `PiThreadFunc` function to collect the partial sums of the rectangle areas within each thread. A copy of this variable will be local to each thread and not require synchronization. Once a thread has completed calculating all the assigned rectangles, the thread should update the global sum with the local partial sum. This update should be protected. Thus, the

protected global update is done once for every thread, rather than once for every rectangle computation.

2. Re-run the updated application through the Thread Profiler to ensure that performance has improved.

## Comparing Performance Runs



1. Click on the "Profiler View" tab and click on the Profile View expansion button.
2. Click on the button (shown at left margin) to remove all data grouping and show the summarized Critical Path data.
3. Drag the original Thread Profiler results of this application from the Tuning Browser to compare with the final results you achieved.
4. Highlight both bars (CTRL+click on bars not highlighted).
5. Examine the results displayed in the "Profile" and Timeline" Views.

Is there a noticeable improvement from the original execution performance?

## Review Questions

---

- Question 1:** Is oversubscription of processor resources by active threads a severe performance problem?
- Yes                      No
- Question 2:** Combining views within the Profile View can give you more information than what is discernible from single views. What can you tell by combining the Objects and Threads views in the Thread Profiler?
- Question 3:** Was the Timeline view useful?
- Question 4:** Can applications instrumented for Timeline view be used in performance comparisons?
- Yes                      No

## Lab. 7: Threaded Programming Methodology

<b>Time Required</b>	Seventy-five minutes
<b>Objectives</b>	<p>In this lab session, you will learn and practice the four steps of Threading Methodology.</p> <p>After successfully completing this lab's activities, you will be able to:</p> <ul style="list-style-type: none"><li>• Apply threading on computational loops within C source code with OpenMP</li><li>• Find actual and potential threading errors, and validate that the errors are fixed using Intel Thread Checker</li><li>• Use Intel Thread Profiler to determine if threaded performance is achieving desired goals, and identify potential bottlenecks to performance</li></ul>





## Activity 1: Compile and Run Serial Code

---

<b>Time Required</b>	Five minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Check that serial code will compile and run; observe behavior of serial application</li> </ul>

The application computes and saves prime numbers between the range of integers given on the command line. The algorithm takes a "brute force" approach and divides each potential prime by possible factors of that number. If one of those factors divides the number evenly, the number is not prime; if all the possible factors have been tried without any being found to divide the number evenly, the number is prime and saved within an array of primes. A running count of the number of primes found is maintained and printed upon completion of the search. A progress update of the percentage of numbers within the given range is printed as the code executes. The progress update value is changed for every 10% completed. The time taken to find the primes will also be printed.

### Build and Run Serial Program

1. Locate and change to the PrimeSingle directory. You should find Visual Studio Solution and project files and a source file, PrimeSingle.cpp, in this directory.
2. Double-click on the solution icon and examine the source file.
3. Be sure the Release configuration is selected and build the executable binary PrimeSingle.exe.
4. After successfully compiling and linking, open a command window, change the directory to the Release folder holding the application, and run from the command line:

```
>> PrimeSingle.exe 1 2000
```

5. Try several different ranges of numbers as command line arguments to the application.

## Activity 2: Analyze Serial Code

<b>Time Required</b>	Ten minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Run serial code through VTune Performance Analyzer to locate portions of the application that use the most computation time. These will be the parts of the code that are the best candidates for threading to make a positive performance impact.</li> </ul>

### Get Baseline Timing

1. Run the serial application to find the primes between 1 and 5,000,000:

```
>> PrimeSingle.exe 1 5000000
```

What time is reported from the application? \_\_\_\_\_ seconds.

### Run VTune Performance Analyzer on Application

1. Switch to the Debug configuration within Visual Studio.
2. Rebuild the application.
3. Start VTune Performance Analyzer and create a new Sampling activity with the debug version of PrimeSingle application to be launched for analysis. Use '1 1000000' as the command line arguments.
4. Run the analysis.

Which function within the code took the most time for execution (most clockticks)?  
\_\_\_\_\_

5. If possible, create a new activity for CallGraph analysis of the same application with the same command line.

What is the average amount of time that each call of the above function takes?  
\_\_\_\_\_usec

## Activity 3: Run OpenMP Code

---

Time Required	Five minutes
Objective	<ul style="list-style-type: none"><li>Build and run OpenMP threaded version of prime finding application</li></ul>

### Build and Run Threaded Program

1. Locate and change to the PrimeOpenMP directory. You should find Visual Studio Solution and project files and a source file, PrimeOpenMP.cpp, in this directory.
2. Using the Debug configuration, build the source file into the executable binary PrimeOpenMP.exe.
3. After successfully compiling and linking, change directories in the command window to the PrimeOpenMP\Debug folder holding the threaded application, and run from the command line:

```
>> PrimeOpenMP.exe 1 5000000
```

What execution time is reported from the threaded application?  
\_\_\_\_\_ seconds.

Compared to the serial application run time, what is the speedup of the threaded version?

[ Speedup = (Serial Time) ÷ (Parallel Time) ] \_\_\_\_\_

## Activity 4: Use Thread Checker to Check for Threading Errors

<b>Time Required</b>	Fifteen minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Use Intel Thread Checker to find the race condition that is causing the code to compute the wrong number of primes</li> </ul>

### Use Thread Checker for Analysis

1. Launch VTune Performance Analyzer on the Windows platform. Create a new Thread Checker activity.
2. Use the debug version of the PrimeOpenMP.exe application to launch. Use '1 200' as the command line arguments. Since the instrumentation from Thread Checker expands the run time of an application, you will not want to run a full test case. Enough numbers to run through the threads will be sufficient for Thread Checker to determine where any threading errors are within the code.
3. Run the application within Thread Checker to generate a diagnostic report.
4. Scan through some of the diagnostics for the run of the application that populate the diagnostic pane in the Thread Checker display. Double-click on one or more of the diagnostic lines to find the source code lines associated with the diagnostic. You can use the tabs at the bottom of the pane to navigate between the diagnostics list and source code.

How many different lines of code have threading errors? \_\_\_\_\_

## Activity 5: Correct Threading Errors

<b>Time Required</b>	Ten minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Correct the errors identified from Thread Checker</li> </ul>

### Correct Errors and Validate Results

1. Within Visual Studio, edit your source code to correct the errors found by Thread Checker.

*Tip:*

Insert critical regions around each use (both read and write) of the affected variables with `#pragma omp critical`.

2. When you have changed the source code, rebuild the application.
3. Re-run the application through Thread Checker.
4. Have all the errors been corrected? If not, go back to Step 1 and repeat until you have successfully handled all the error diagnostics reported by Thread Checker.

### Build and Run Threaded Program

1. Once the threading errors have been removed from the program, re-run the application within the command window:

```
>> PrimeOpenMP.exe 1 5000000
```

What execution time is reported from the threaded application?  
\_\_\_\_\_ seconds.

Compared to the serial application run time, what is the speedup of the corrected threaded version?

[ Speedup = (Serial Time) ÷ (Parallel Time) ] \_\_\_\_\_

## Activity 6: Find Threading Performance Issues

<b>Time Required</b>	Ten minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Use Intel Thread Profiler to find any problems that might be hampering the parallel performance</li> </ul>

The prime finding application needs to be rebuilt with the performance measurement source code instrumentation flag in order to identify potential performance bottlenecks resulting from how threads interact with each other. This will allow either OpenMP or explicit threads analysis can be used within Intel Thread Profiler.

### Instrument and Run Threaded Program

1. Open the PrimeOpenMP project properties dialog box. Within the C/C++ folder, on the Command Line, change the /Qopenmp flag to /Qopenmp\_profile. This is the source code instrumentation flag.
2. Rebuild the application with the new compile flag.
3. Create a new Thread Profiler activity. Be sure to choose the "Threaded (Windows\* API or POSIX Threads)" radio button in the Threading Type selector. Run the executable with a full range of numbers (1 to 1000000 should be sufficient). A "production-sized" data set is used since this would be the data size normally used. Performance problems may not be evident if the data set used for analysis is smaller than normal.

### Use Thread Profiler for Analysis

1. Once the analysis is available, click through the tabs to examine the different views.

From the Timeline view, can you get to the transition source that is showing some thread inactivity?

## Activity 7: Reduce the Number of printf Calls

<b>Time Required</b>	Five minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Correct the problem of implicit synchronization issue from calling printf more times than necessary</li> </ul>

### Modify Code to Control Number of Progress Updates

1. Edit the PrimeOpenMP.cpp source code to only print the minimum required progress update output statements and, consequently, remove the implicit synchronization from unnecessary use of thread-safe library calls.

```
void ShowProgress( int val, int range )
{
    int percentDone;

    static int lastPercentDone = 0;

    #pragma omp critical
    {
        gProgress++;

        percentDone = (int)((float)gProgress/
        (float)range*200.0f+0.5f);

        if( percentDone % 10 == 0 && lastPercentDone <
        percentDone / 10){

            printf("\b\b\b\b\b%3d%%", percentDone);

            lastPercentDone++;

        }
    }
}
```

2. Compile the source file into the executable binary PrimeOpenMP.exe. (Be sure the /Qopenmp\_profile flag has been changed back to just /Qopenmp for compilation.) Within the command window, run the executable with an appropriate range of numbers to test.

```
>> ./PrimeOpenMP 1 5000000
```

What execution time is reported from the load balanced application?  
 \_\_\_\_\_ seconds.

Compared to the serial application run time, what is the speedup of the current threaded version?

[ Speedup = (Serial Time) ÷ (Parallel Time) ] \_\_\_\_\_

## Update Serial Code

1. Modify the serial version of the code to use the same ShowProgress function that the OpenMP code uses. Recompile and run the serial application.

What execution time is reported from the improved serial application?  
 \_\_\_\_\_ seconds.

Compared to the updated serial run time, what is the speedup of the current threaded version?

[ Speedup = (New Serial Time) ÷ (Parallel Time) ] \_\_\_\_\_



## Activity 8: Modify Critical Regions

<b>Time Required</b>	Ten minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Modify the explicit critical regions to reduce contention and limit the amount of time threads spend in each region</li> </ul>

### Use *InterlockedIncrement* and Local Variables

- For each critical region within the PrimeOpenMP.cpp source file, modify the code to replace the OpenMP critical pragmas by using the Windows InterlockedIncrement function.

First modification:

```
#pragma omp parallel for
    for( int i = start; i <= end; i+= 2 ){
        if( TestForPrime(i) ) {

            globalPrimes[InterlockedIncrement(&gPrimesFound)] = i;

        }

        ShowProgress(i, range);
    }
```

Second modification:

```
    long percentDone, localProgress;

    static int lastPercentDone = 0;

    localProgress =
        InterlockedIncrement(&gProgress);

    percentDone = (int)((float)localProgress/
        (float)range*200.0f+0.5f);
```

- Rebuild the source file into the executable binary PrimeOpenMP.exe. Run the application within the command window with an appropriate range of numbers to test.

```
>> ./PrimeOpenMP 1 5000000
```

What execution time is reported from the threaded application?  
\_\_\_\_\_ seconds.

Compared to the (new) serial application run time, what is the speedup of the threaded application?

[ Speedup = (Serial Time) ÷ (Parallel Time) ] \_\_\_\_\_

## Activity 9: Fix Load Imbalance

<b>Time Required</b>	Five minutes
<b>Objective</b>	<ul style="list-style-type: none"> <li>Correct the load balance issue between threads by distributing loop iterations in a better fashion</li> </ul>

### Use OpenMP schedule Clause

1. Within Visual Studio, edit your source code to remedy the load imbalance problem identified by Thread Profiler. To do this, add a schedule clause to the OpenMP parallel region pragma.

```
#pragma omp parallel for schedule(static, 8)
for( int i = start; i <= end; i += 2 )
{
    if( TestForPrime(i) )

        globalPrimes[InterlockedIncrement(&gPrimesFound)] = i;

        ShowProgress(i, range);
}
```

2. Compile the source file into the executable binary PrimeOpenMP.exe and run the executable with an appropriate range of numbers to test.

```
>> PrimeOpenMP.exe 1 5000000
```

What execution time is reported from the load balanced application?  
 \_\_\_\_\_ seconds.


3. Build the Release configuration of the threaded application and run.

What is the execution time? \_\_\_\_\_ seconds.

What execution time is reported from the Release build application?  
 \_\_\_\_\_ seconds.

Compared to the serial application run time, what is the final speedup of the threaded version?

[ Speedup = (Serial Time) ÷ (Parallel Time) ] \_\_\_\_\_



## **Lab. 8: Scalability of Threaded Applications**



## **Workload-dependent Scaling - An Overview**

---

Activity 1, comprising four parts, discusses an example of workload dependent scaling with a multi-threaded application. Scaling of multithreaded applications may vary based on workloads. Workloads with balanced data set division among the threads and with larger data set may scale well. On the other hand, workloads with imbalanced division or smaller data sets may show poor scaling. Studying workload statistics is one of the important areas to consider when multi-threading an application.

While root causing poor scaling of a multithreaded application, it is recommended that different sizes of workloads be run to determine whether scaling performance changes with a different workload type/size. Another recommended step is to use Intel® Thread Profiler in identifying serial vs. parallel time.

In this four-part activity, you will examine a sample application which shows a difference in multithreaded performance improvement due to the difference in size of workloads.

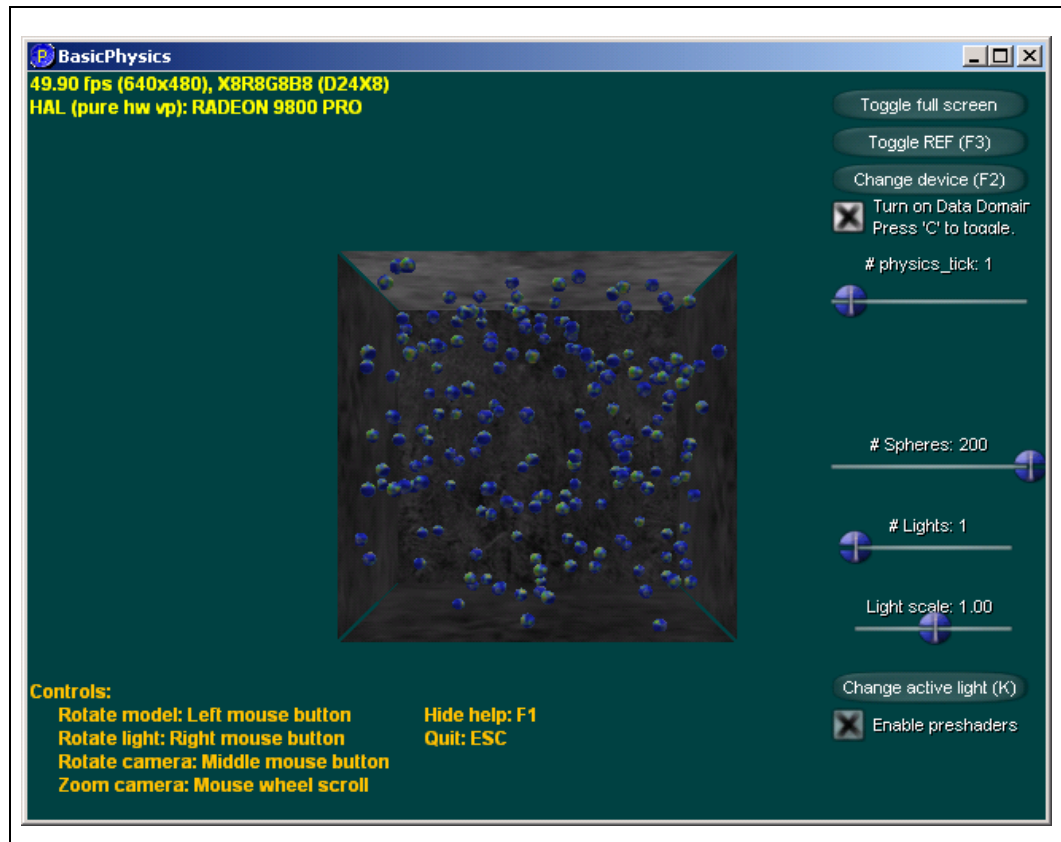
### **Introduction to Sample Application**

The sample application used here is a DirectX application based on the BasicHLSL sample included in DX SDK.

This example demonstrates a number of spheres in a box with an implementation of collision detection mechanisms for sphere-sphere and sphere-plane based on Newtonian physics. The GUI of the BasicPhysics demo (see [Figure 8.1](#)) provides sliders to add or remove spheres from the scene and change the number of times the physics computations are done before actually rendering.

The GUI displays the frame rate at the top left corner. Besides the GUI, there is a checkbox to turn on threading for data decomposition. When unchecked, all computation and rendering is done by a single thread; when checked, the rendering is done on the main thread and the position computations for the sphere objects are divided among two worker threads.

**Figure 8.1. BasicPhysics Application's User Interface**



While multi-threading an application, it is important to study characteristics of typical workloads. If the workload size is too small, multi-threading may cause an extra overhead. Another factor is to maintain a good balance of work among all threads.

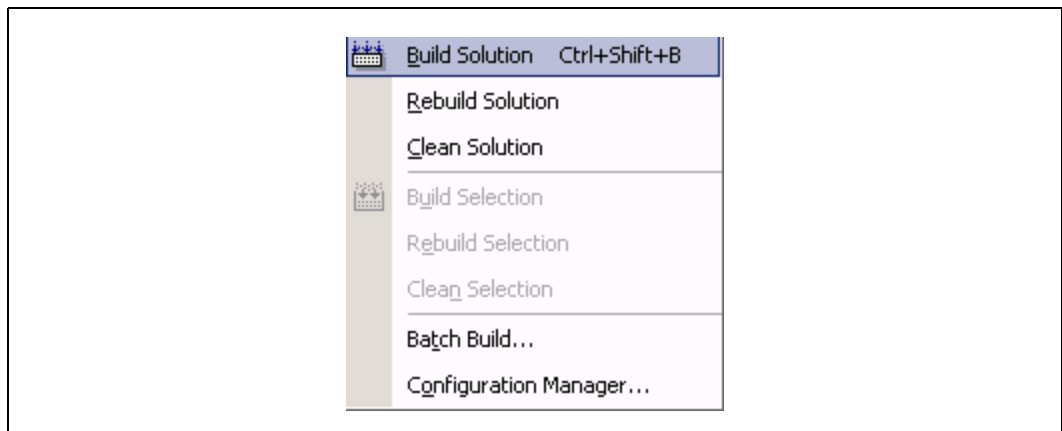


## Activity 1a: Build the Application

---

1. Locate the 'BasicPhysics.sln' solution file in the Scalability\Workload Scaling Lab directory.
2. Double-click the icon to launch Visual Studio .NET 2005 with this solution loaded.
3. Use Ctrl+Shift+B or the Build menu to build the application after the .sln file is opened.

**Figure 8.2. The Build Menu**



4. Run the application using Ctrl + F5.

**Note:** The application may take some time to launch, as it loads and initializes meshes for all spheres.

## Activity 1b: Characterize Multithreaded Implementation with Large Data Set

---

1. Set 'physics\_tick' = 10 and '# Spheres' = 200 by moving sliders to right.
2. Uncheck 'Turn On Data Domain Decomposition' box. Once it has stabilized, note the frame rate (in space below) displayed at the top left corner of the window.
3. Check 'Turn On Data Domain Decomposition' box and note frame rate displayed at the top left corner of the window.
4. Compute Multi-Threaded scaling of the application with the large data set.

ST Frame Rate with 200 spheres and physics\_tick 10: \_\_\_\_\_FPS

MT Frame Rate with 200 spheres and physics\_tick 10: \_\_\_\_\_FPS

MT Speedup: \_\_\_\_\_


## Activity 1c: Examine Thread Performance with Large Data Set

This activity should be run while the multithreaded version of the BasicPhysics application is running with 200 spheres.

1. From a command line window, launch the PerfMon performance monitor while the BasicPhysics application is running.
2. Click on the "Add" button (or use Ctrl-I).
3. Within the "Add Counters" window, select "Threads" from the Performance Object pull-down menu.
4. Be sure the "% Processor Time" counter is selected and add each of the BasicPhysics threads from the Instances pane. There should be three threads from BasicPhysics.
5. Click the Close button on the "Add Counters" window.
6. Remove the other counters that were active when PerfMon was started in order to focus on the BasicPhysics threads.

After some execution time has passed, you should see something similar to the graphing pane excerpt shown in [Figure 8.3](#).

**Figure 8.3. Threads time with 200 spheres and physics\_tick 10**



## Activity 1d: Characterize Multithreaded Implementation with Small Data Set

1. While still running the BasicPhysics application, bring the GUI window to the front.
2. Set 'physics\_tick' = 5 and '# Spheres' = 100 by moving sliders to right.
3. Uncheck 'Turn On Data Domain Decomposition' box. Once it has stabilized, note the frame rate (in space below) displayed at the top left corner of the window.
4. Check 'Turn On Data Domain Decomposition' box and note frame rate displayed at the top left corner of the window.
5. Compute Multi-Threaded scaling of the application for 100 spheres.


ST Frame Rate with 100 spheres and physics\_tick 5: \_\_\_\_\_ FPS

MT Frame Rate with 100 spheres and physics\_tick 5: \_\_\_\_\_ FPS

MT Speedup: \_\_\_\_\_

6. While running the BasicPhysics application in multithreaded mode, bring the PerfMon window to the front.  
You should see something similar to the graphing pane excerpt shown in [Figure 8.4](#).

**Figure 8.4. Threads time with 100 spheres and physics\_tick 5**



### **Lab. 8: Scalability of Threaded Applications**

Comparing your calculated speedup results from Activities 1c and 1d, it should be clear that better scaling is possible when time spent in the parallel (threaded) execution is much larger than time spent in serial section. PerfMon results in Figure 2 (workload of 200 spheres, physics\_tick=10) and Figure 3 (workload of 100 spheres, physics\_tick=5) show the relationship of parallel to serial time during execution.

## Activity 2: Illustrate Effects of Load Balance in Multithreaded Implementation

---

1. While still running the BasicPhysics application, bring the GUI window to the front.
2. Set 'physics\_tick' = 10 and '# Spheres' = 200 by moving sliders to right. Also, ensure that the 'Turn On Data Domain Decomposition' box is checked.
3. The "Balance pct." slider should have been pre-set to 50. This divides the position computations evenly between the two worker threads. Move the slider to the left to a position close to 33. This will assign twice as many objects to the second thread than are assigned to the first thread.
4. Note the frame rate below:

MT Frame Rate with 200 spheres, physics\_tick 10, and Balance pct. 30:  
 \_\_\_\_\_FPS

**Question 1:** Is there any difference in the graphs of the two worker threads within the PerfMon graphing pane?

5. Move the "Balance pct." Slider to the left to a position close to 10. Computation for nine times as many objects will be assigned to the second thread in this case.
6. Note the frame rate below:

MT Frame Rate with 200 spheres, physics\_tick 10, and Balance pct. 10:  
 \_\_\_\_\_FPS

**Question 2:** How do the graphs of the two worker threads differ under this load assignment as shown by PerfMon?

**Question 3:** How do the threaded frame rates of the imbalanced work loads compare to the serial and balanced multithreaded workloads?

## Activity 3: Measuring Synchronization Object Overhead

<b>Objectives</b>	<ul style="list-style-type: none"><li>• Measure the relative speed of different synchronization objects in an uncontended fashion</li></ul>
-------------------	---

1. Locate the Scalability/SyncObjectBenchmark folder and double-click on the Visual Studio solution icon. Examine the code to understand what is being timed in the application. Rebuild the project.
2. From the Debug menu, run the executable by choosing the "Start Without Debugging" command.
3. Fill in the relevant entries in the table provided below, from the output generated.

File	Time	Relative Speed
InterlockedIncrement		<b>1.0X</b>
Critical Section		
Critical Section with Spincount		
Mutex		
Semaphore		

**Question 1:** Is there any difference in the relative speeds if the number of loops, each synchronization object is tested on, is changed? Why or why not?

**Question 2:** Does a different spin count change the relative speed of the use of the Critical Section object?

## Activity 4: Measuring Front-side Bus Saturation

<b>Objectives</b>	<ul style="list-style-type: none"> <li>• Observe impact of FSB saturation on scalability using Stream benchmark</li> <li>• Learn use of appropriate VTune performance event to monitor bus utilization</li> </ul>
-------------------	---

This lab is designed for a dual socket multi-core platform with dual independent buses. You will build and run the Stream benchmark application on multiple cores. Stream will report execution time and a "Copy Function Rate" in MB/s. We will compare this latter measure with VTune event data that is used to compute bus utilization.

1. Locate the Scalability/Front-side Bus folder and double-click on the Visual Studio solution icon.
2. Re-build ScalabilityLab-Bus project.  
(The binary is made available at the Release directory.)
3. Launch Windows Task Manager and bring up the "Performance" tab.
4. Open a command window and change directory to the Scalability/Front-side Bus/Release directory.
5. Run "scalabilitybus0.bat" at the command window.  
This runs Stream on Core 0 (Socket 0).
6. Note the "Copy Function" Rate and Run Time in the table provided below.
7. Run "scalabilitybus0123.bat" at the command window.  
This runs 4 Stream processes (Cores 0, 1, 2, 3 used).
8. Note the "Copy Function" Rate and Run Time in the table provided below.
9. If your system has 8 cores available, you can run the "scalabilitybus0246.bat" and the "scalabilitybus0-7.bat" command files. Note the "Copy Function" Rate and run time for each in the table provided below.

File	Copy Function Rate	Run Time
scalabilitybus0.bat		
scalabilitybus0123.bat		
scalabilitybus0246.bat		
scalabilitybus0-7.bat		

10. Start VTune and set up a new activity to track the "Bus Clockticks" (CPU\_CLK\_UNHALTED.BUS) and "Bus Data Ready from the Processor" (BUS\_DRDY\_CLOCKS.THIS\_AGENT) events and execute the "scalabilitybus0.bat" command file.



**Lab. 8: Scalability of Threaded Applications**

11. Compute bus bandwidth for the core and %bus utilization. Note the results in the table below.
12. Drill down to source to note high clock tick areas.
13. Run the VTune activity with "scalabilitybus0123.bat" (as well as "scalabilitybus0246.bat" and "scalabilitybus0-7.bat", if possible).
14. Compute total bus bandwidth and %bus utilization for each run. Note the results in the table provided below.

File	Total Bus Bandwidth	% bus utilization
scalabilitybus0.bat		
scalabilitybus0123.bat		
scalabilitybus0246.bat		
scalabilitybus0-7.bat		

## Activity 5: Identifying False Sharing in Threaded Applications

<b>Objectives</b>	<ul style="list-style-type: none"> <li>• Observe impact of false sharing on scalability</li> <li>• Learn use of appropriate VTune performance event to compare and contrast false sharing vs. no false sharing</li> </ul>
-------------------	---

This lab is designed for a dual socket multi-core platform with dual independent buses. You will build and run the Stream benchmark application on multiple cores. Stream will report execution time and a "Copy Function Rate" in MB/s. We will compare this latter measure with VTune event data that is used to compute bus utilization.

1. Locate the Scalability/False sharing folder and double-click on the Visual Studio solution icon. Re-build ScalabilityLab-FS project. The binary will be available in Release directory.
2. Open a command window and change directory to the Scalability/False sharing/Release directory.
3. Run "scalabilityfs-1t-fs.bat".  
This runs code with false sharing using 1 thread fixed to Core0 (Socket 0).
4. Note run time reported at table provided below.
5. Run "scalabilityfs-4t-fs.bat".  
This runs code with false sharing using 4 threads with one thread fixed to each core.
6. Note run time reported at table provided below.
7. Use the single core run time to compute speedup.
8. Run "scalabilityfs-4t-nofs.bat".  
This runs code with no false sharing using 4 threads with one thread fixed to each core.
9. Note run time reported at table provided below.
10. Use the single core run time to compute speedup.

File	Run Time	Speedup
scalabilityfs-1t-fs.bat		<b>1.0X</b>
scalabilityfs-4t-fs.bat		
scalabilityfs-4t-nofs.bat		

**Lab. 8: Scalability of Threaded Applications**

11. Launch the VTune Performance Analyzer. Create an activity to include the following events:

- Clockticks ..... CPU\_CLK\_UNHALTED.CORE
- Instructions Retired ..... INST\_RETIRED.ANY
- Memory Order Machine Clear ..... MACHINE\_NUKES.MEM\_ORDER
- 2nd Level Cache Read Misses ..... MEM\_LOAD\_RETIRED.L2\_MISS
- Bus Data Ready From the Processor ..... BUS\_DRDY\_CLOCKS.THIS\_AGENT

12. Run following batch files within VTune using the events listed above. Use a new event for each run. Fill in the table provided below from the total event counts for each run. Drill down to source view in each and compare locations of high numbers of event activities.

- scalabilityfs-1t-fs.bat
- scalabilityfs-4t-fs.bat
- scalabilityfs-4t-nofs.bat

	Sum of events on all 4 cores		
	1T-FS	4T-FS	4T-NOFS
Clockticks			
Instructions Retired			
Memory Order Machine Clear			
2nd Level Cache Read Misses			
Bus Data Ready From the Processor			


## Review Questions

---

- Question 1:** How does the scaling in Activity 1b and 1d compare?
- Question 2:** How is multithreaded scaling affected by workload size?
- Question 3:** How would you identify the amount of serial and parallel time for a workload?
- Question 4:** What are some of the initial steps in root causing poor MT scaling with a particular workload?

## Answers

1. Activity 1 demonstrates higher potential for scaling, as workload size is larger, hence more time spent in parallel sections.
  2. In case of smaller size workloads, effective work done by threads may be less as compared to larger data set. Also if the data set is too small, overhead for maintaining threads may potentially be higher than gain achieved by multi-threading.
  3. Run Perfmon/Intel® Thread Profiler monitoring %processor time, %user time and elapsed time. This will give an estimate of serial vs. parallel time for that workload. (Parallel time: When multiple threads are active concurrently. Serial time: When only 1 thread is active at a given time).
  4. Collect Intel® Thread Profiler/Perfmon data to understand serial vs. parallel time spent in workload.
- Characterize various workloads with same application to examine if MT scaling changes with workloads.
- Analyze MT implementation to identify any bottlenecks (tools like Intel® Thread Profiler may be helpful).



# Appendices



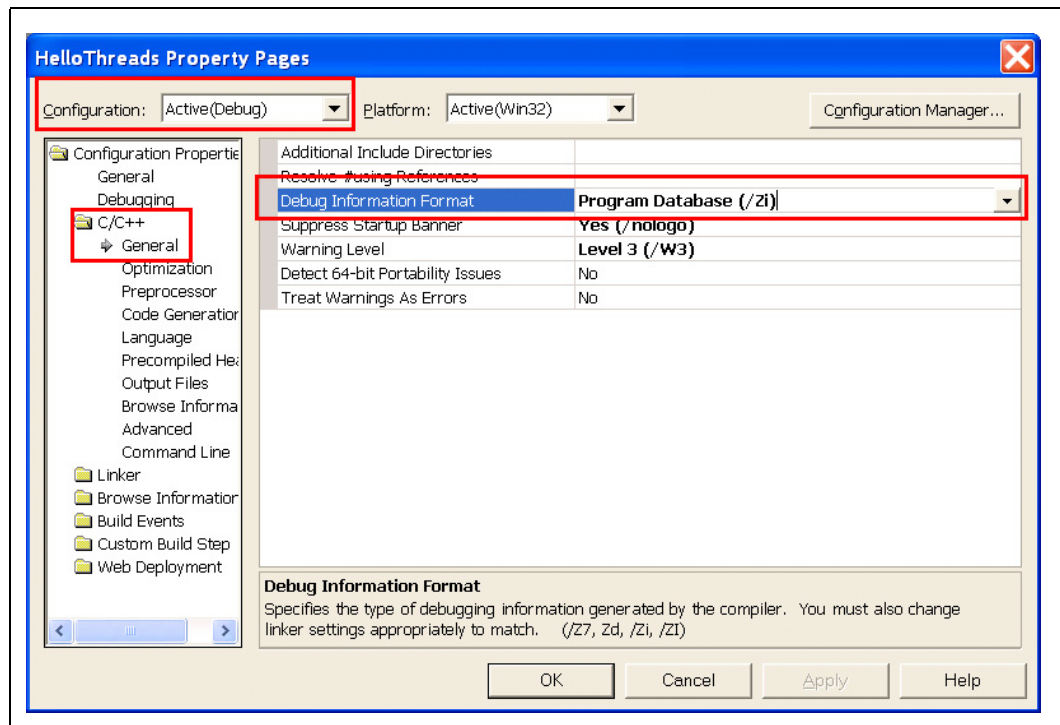
## Appendix A: Setting-up Intel® Thread Checker

### Setting Compile & Link Options

In order to compile applications for use with Thread Checker, certain debug features, optimization levels, and system library choices need to be made. This section contains the steps that detail the setting that should be used and how to check and set them within a Microsoft Visual Studio project configuration.

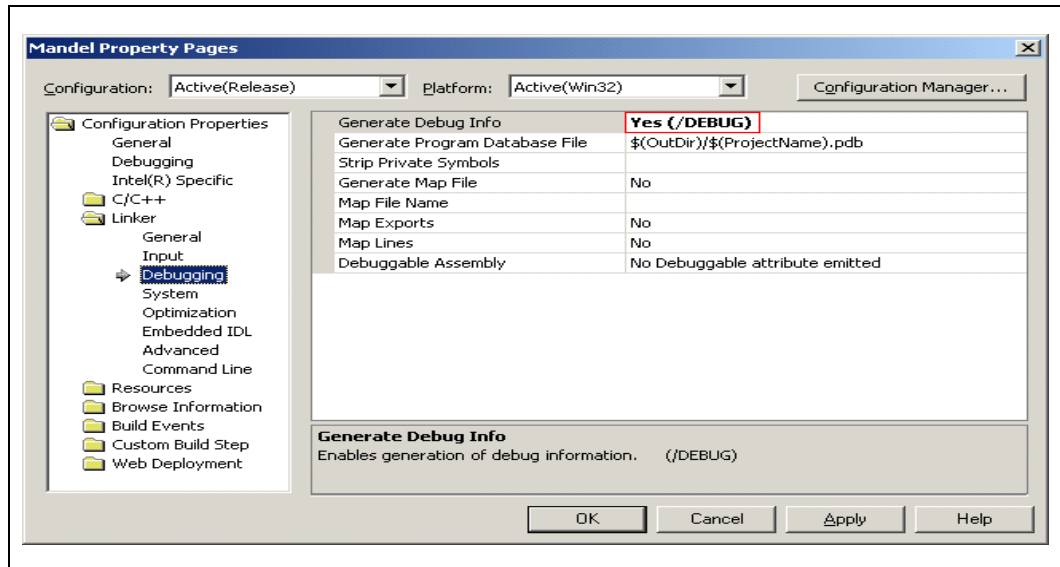
1. Ensure that the Debug options are selected, as shown in [Figure A-1](#).

**Figure A-1. Project Setting – C/C++ Folder – Debug Options**



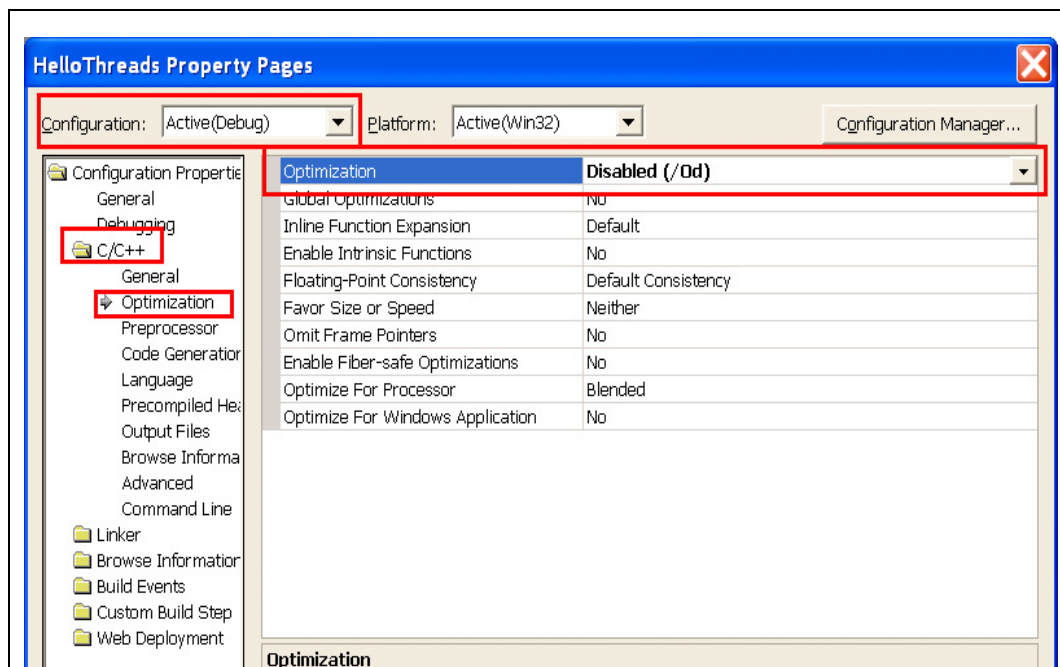
2. Ensure that the Debug symbols are preserved during the link phase, as shown in Figure A-2.

**Figure A-2. Linker Settings – Debugging Folder**



3. Ensure that the Optimization is disabled, as shown in Figure A-3.

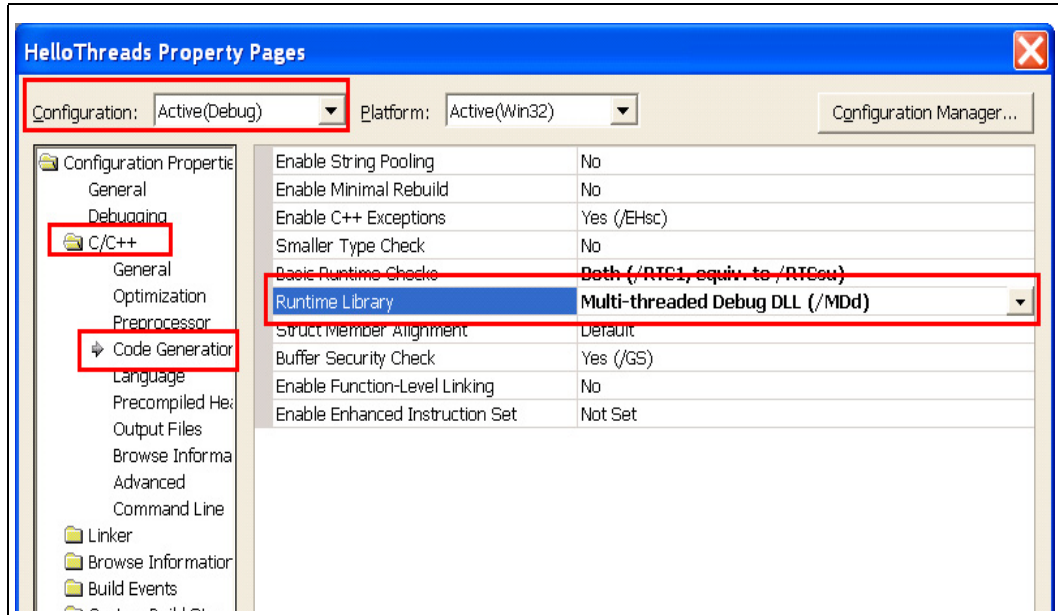
**Figure A-3. Project Settings – C/C++ Folder – Optimization Options**





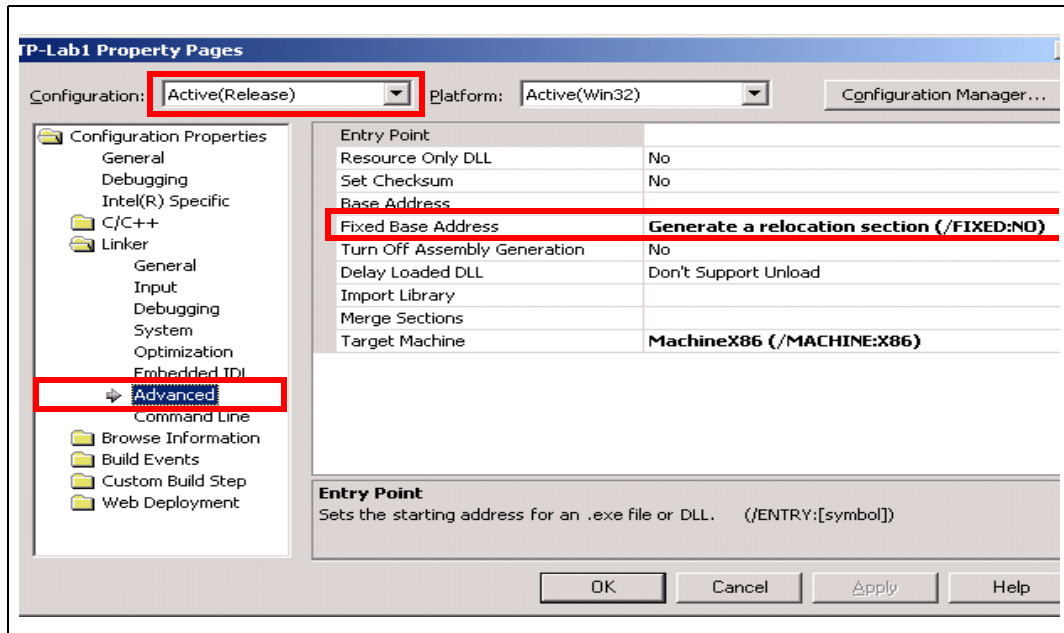
4. Ensure that the thread-safe libraries are selected, as shown in Figure A-4.

Figure A-4. Project Settings – C/C++ Folder – Thread-safe Libraries Options



5. Ensure that the application is built with the /fixed:no option from the Linker->Advanced attribute as shown in Figure A-4.

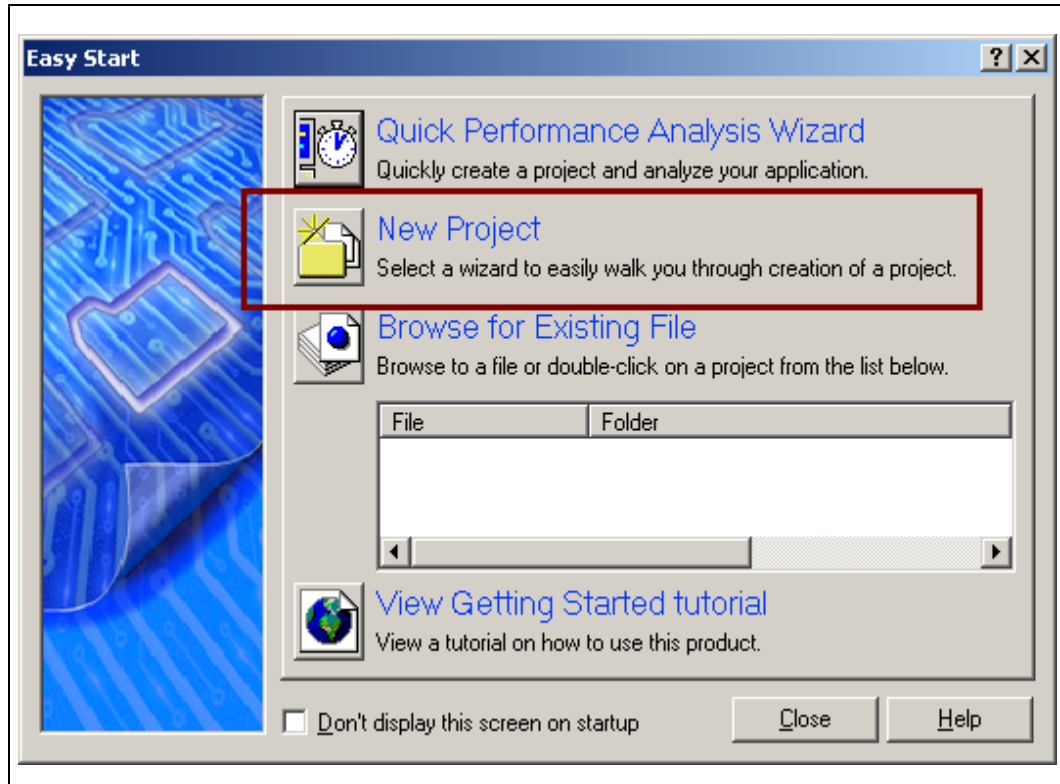
Figure A-5. Linker Settings – Advanced Attributes



## Creating a New Project and Thread Checker Activity

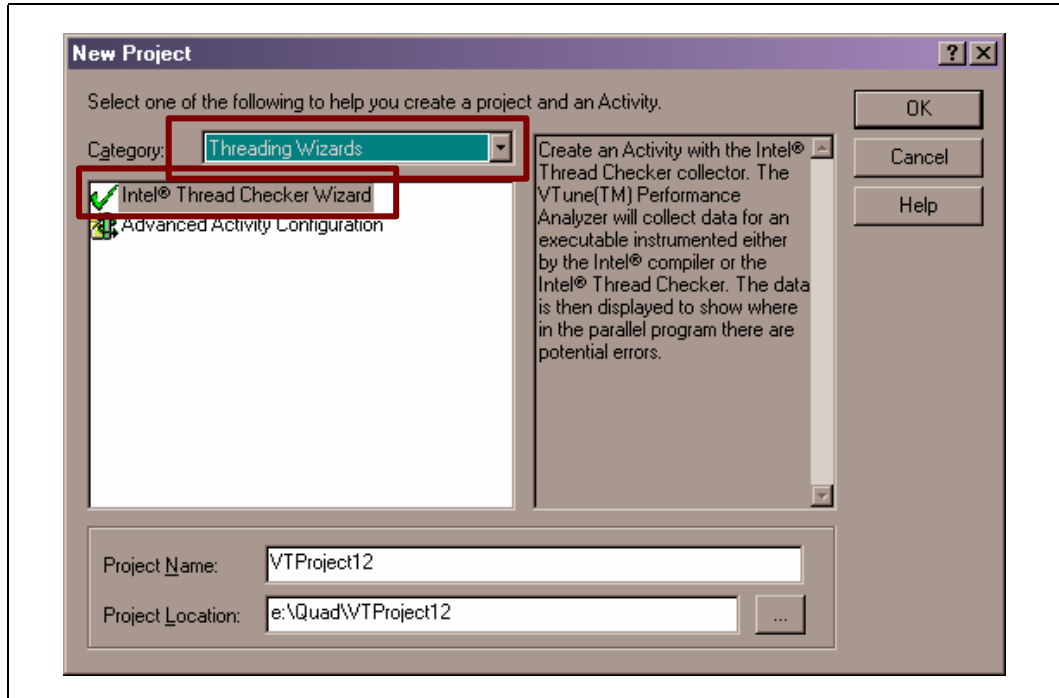
1. Start VTune Performance Analyzer and select New Project, as shown in Figure A-6.

**Figure A-6. VTune™ Analyzer Easy Start Menu**



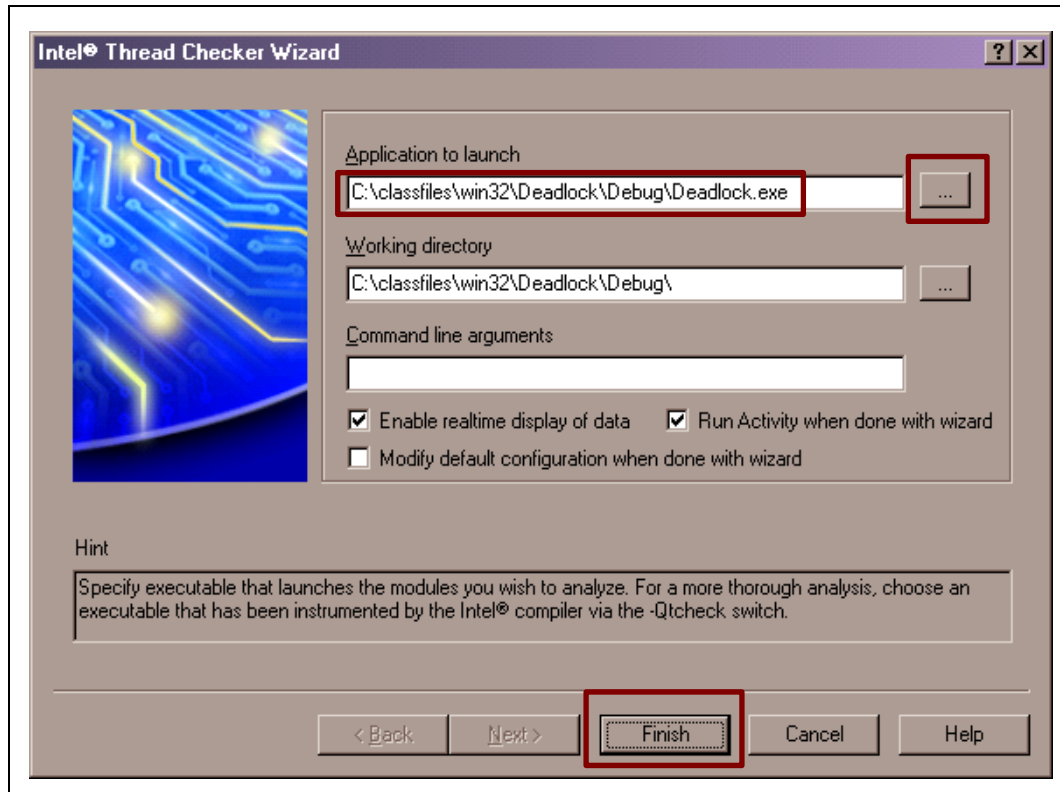
2. From the Category Threading Wizards, select the Intel® Thread Checker Wizard, as shown in [Figure A-7](#).

**Figure A-7. Intel® Thread Checker New Project Menu**



3. Set up the application to be run by using the browse “...” button (highlighted in red in Figure A-8.) and then click Finish to start Intel Thread Checker.

**Figure A-8. Intel® Thread Checker Wizard**



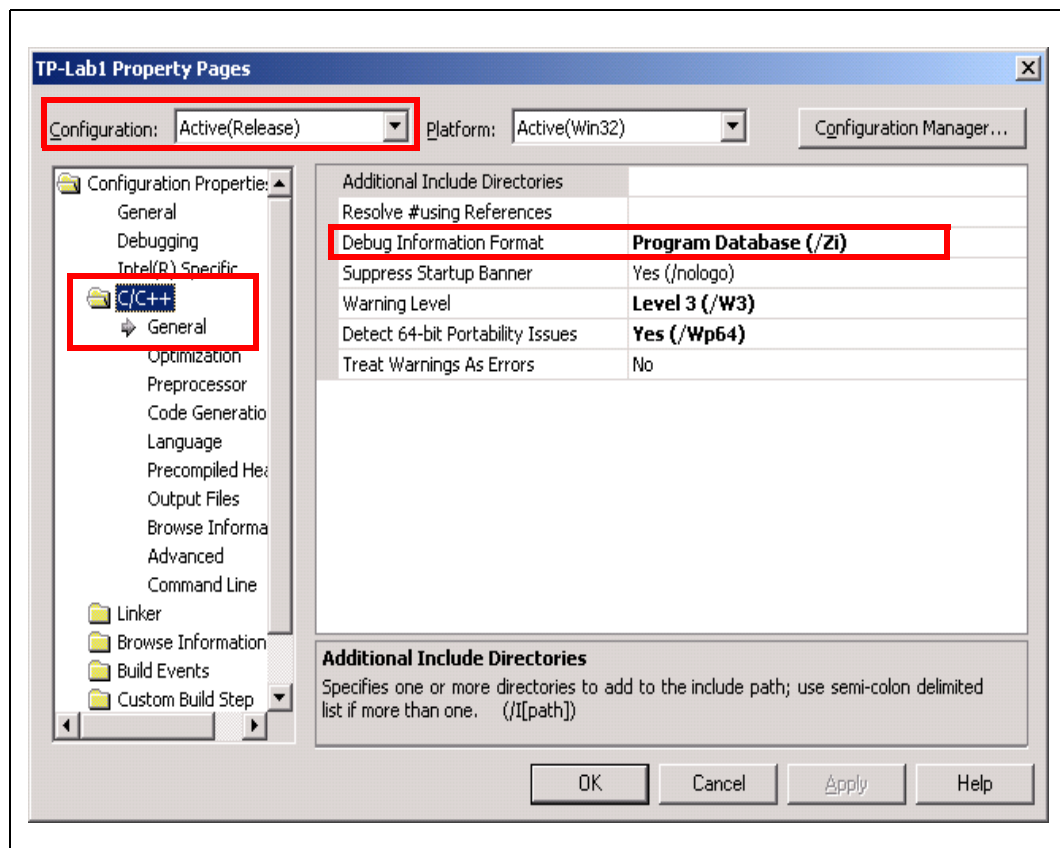
## Appendix B: Setting-up Intel® Thread Profiler

### Setting Compile & Link Options

In order to compile applications for use with Intel® Thread Profiler, certain debug features and system library choices need to be made. The steps in this section detail the setting that should be used and how to check and set them within a Microsoft Visual Studio project configuration.

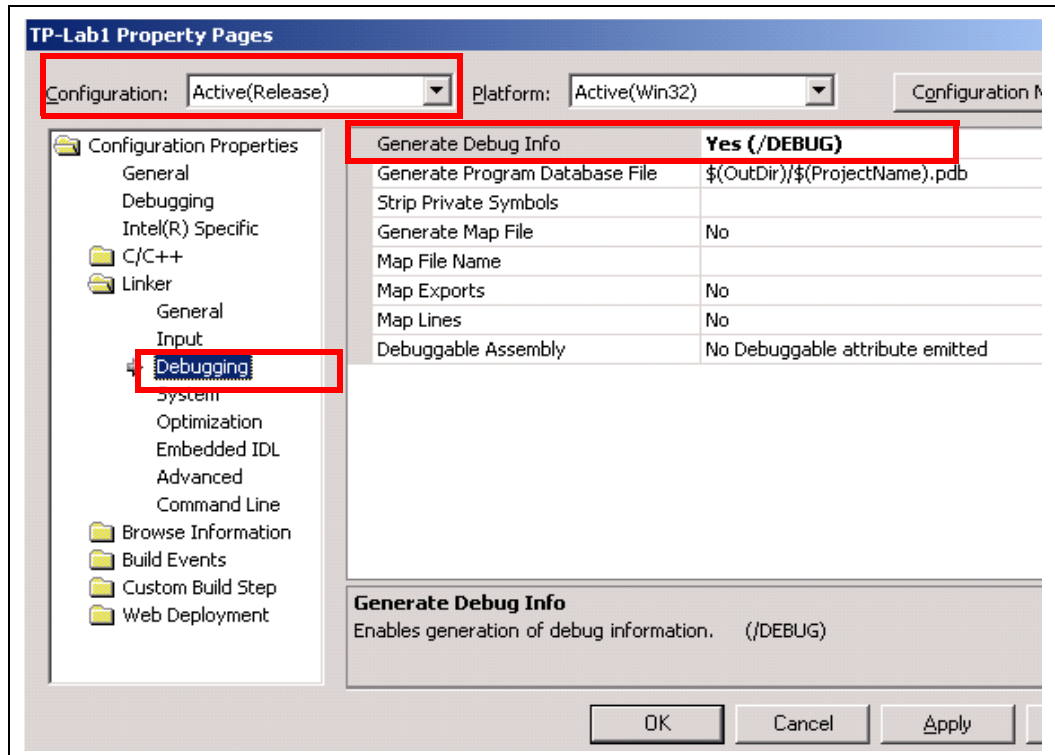
1. From the Build menu, select Configuration Manager... and then select the Release build.
2. Make sure that Debug format is specified as shown in Figure B-1.

**Figure B-1. Specifying Debug Format**



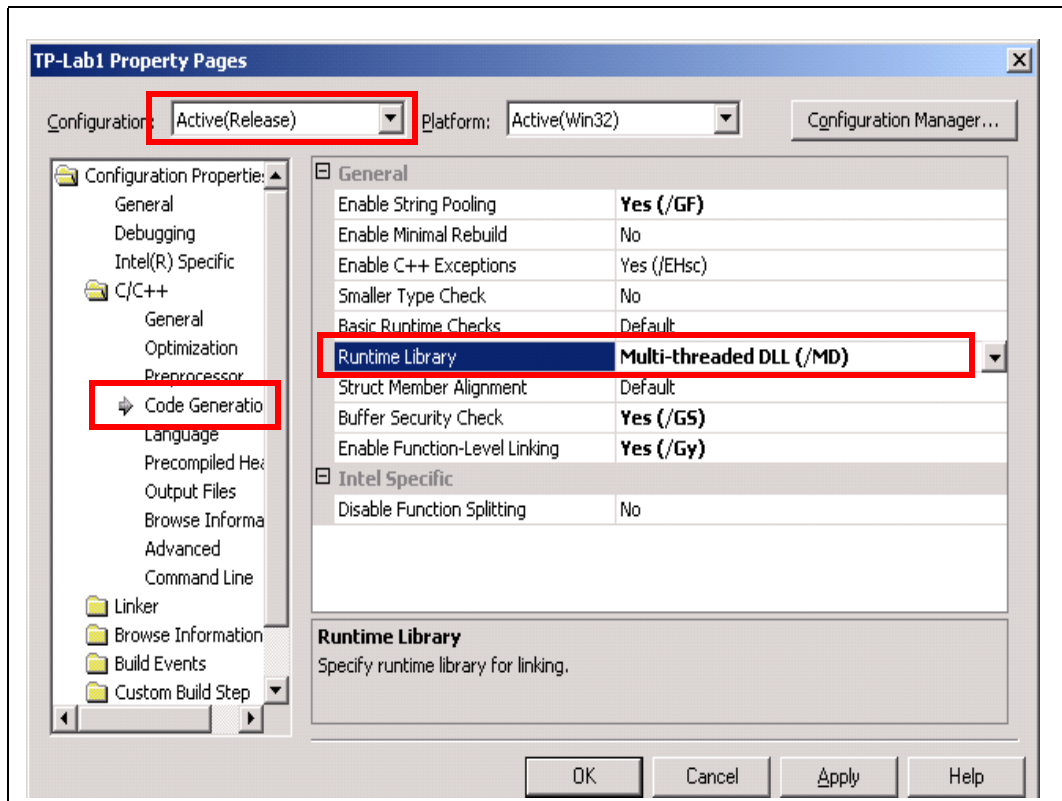
3. Make sure that the debug symbols are generated for the application as shown in Figure B-2.

**Figure B-2. Generating Debug Symbols**



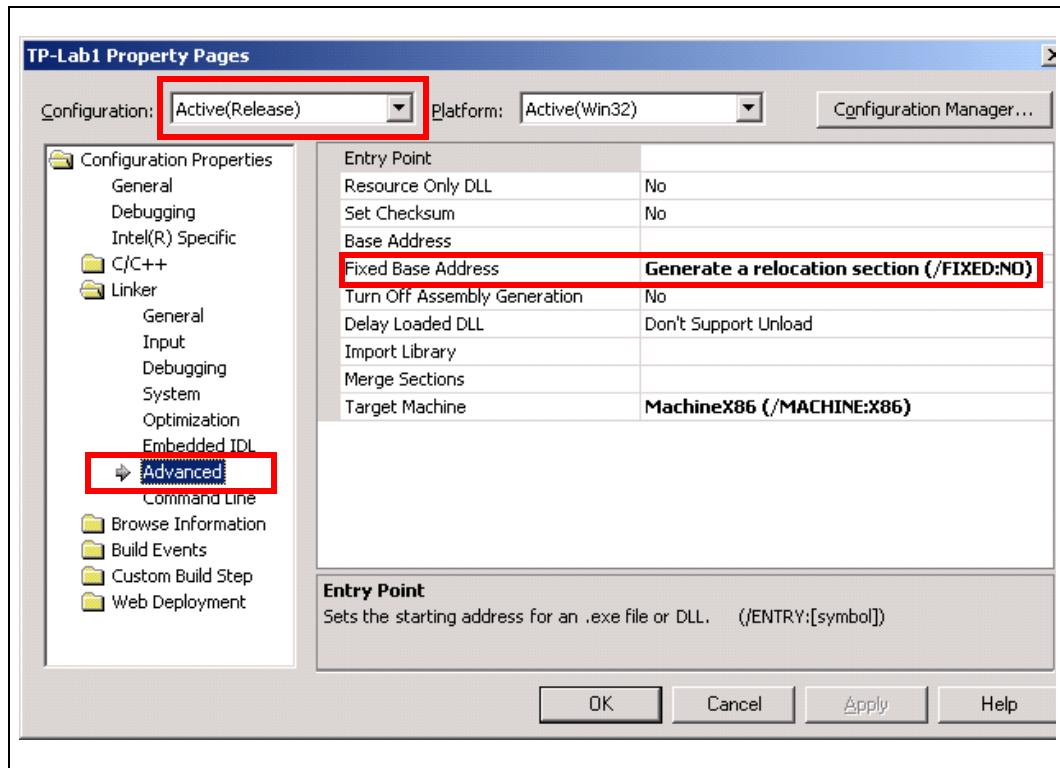
4. Make sure that thread safe libraries have been selected as shown in Figure B-3.

**Figure B-3. Selecting Thread Safe Libraries**



5. Make sure that the application is built with the **/fixed:no** option from the Linker->Advanced attribute as shown in Figure B-4.

**Figure B-4. Specifying Relocatable Binary**



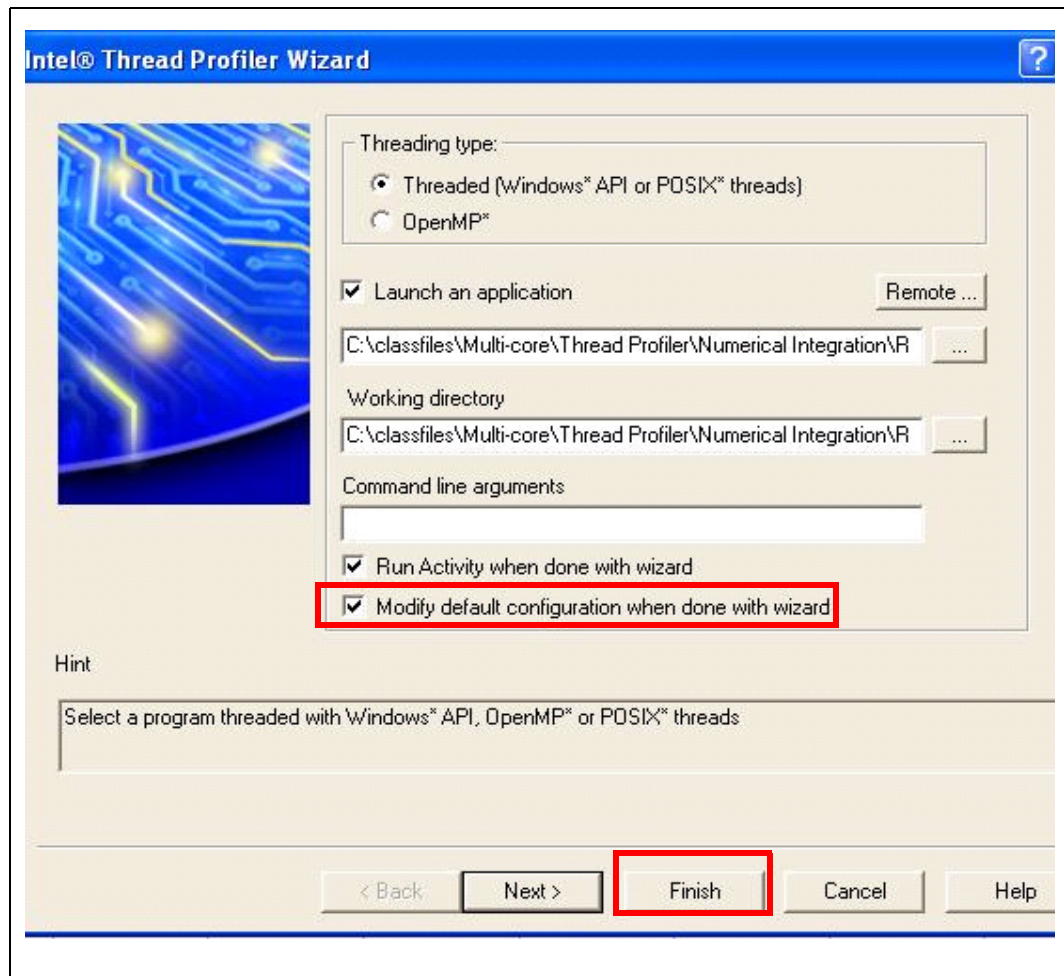
6. From the Build menu, select Build Solution to build. The application will now be ready to be run under the Thread Profiler.



## Creating a New Project and Thread Checker Activity

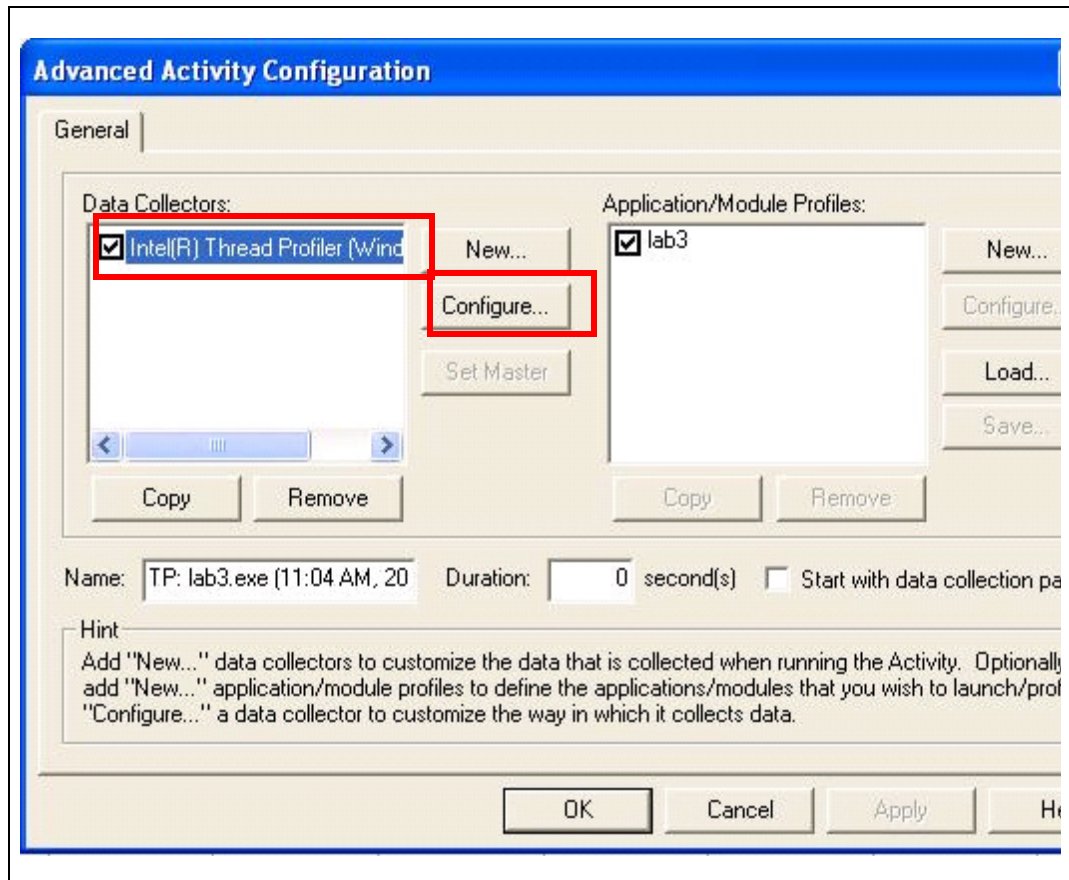
1. Start VTune™ Performance Analyzer and select New Project.
2. From the Category Threading Wizards, select the Intel® Thread Profiler Wizard. (Figure B-5)
3. Within the Thread Profiler Wizard, be sure to click on the "Threaded (Windows\* API or POSIX\* threads)" radio button is selected.
4. Choose an application to be analyzed and enter any command line arguments required.

Figure B-5. Thread Profiler Wizard



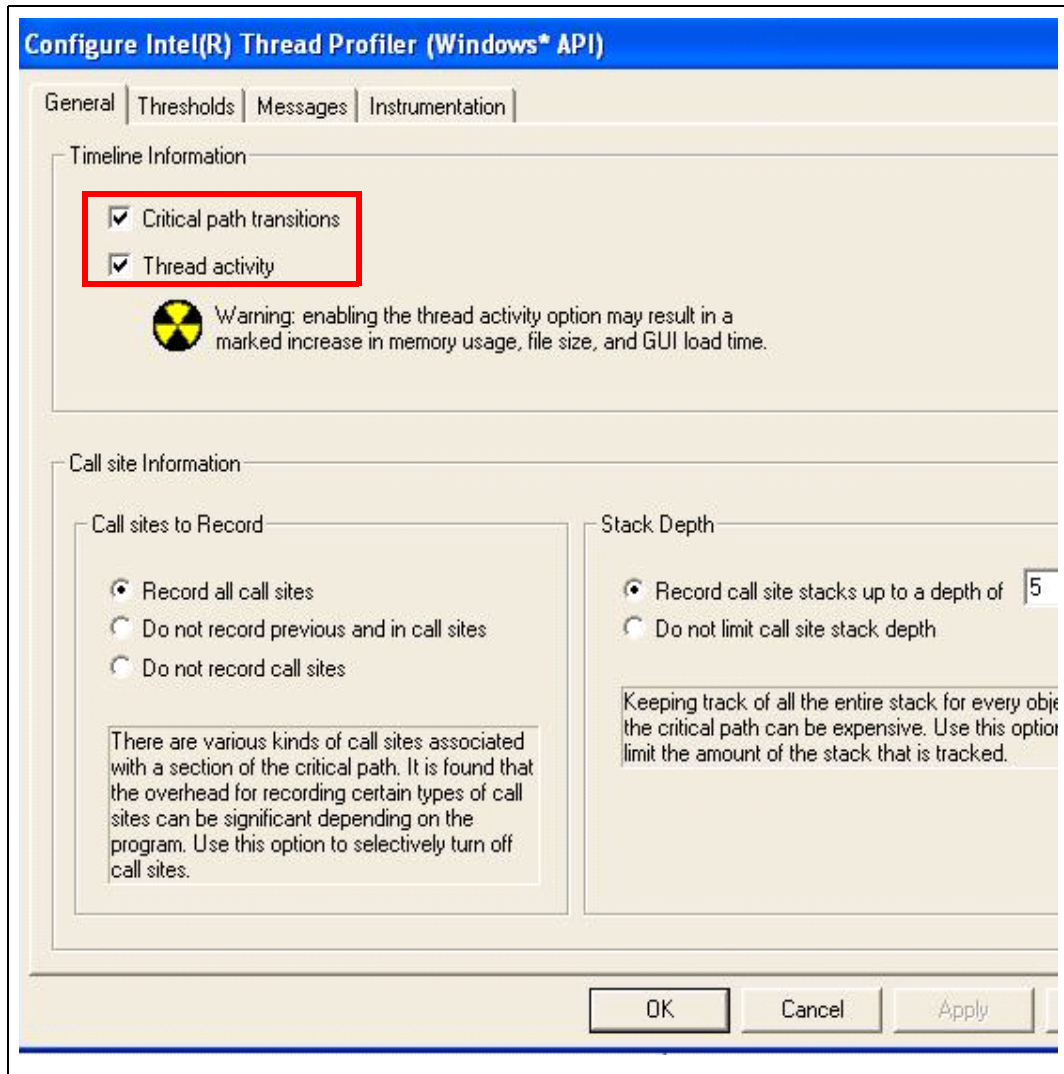
5. Select "Thread Profiler" and click on "Configure" (Figure B-6). In the resulting pop-up dialog, select the "Miscellaneous" Tab.

**Figure B-6. Advanced Activity Configuration**



6. Select the "Thread Activity" and "Transitions" check boxes (Figure B-7). This selection will issue a warning that the Thread Profiler may take longer to run. Accept this warning by clicking on "Yes".

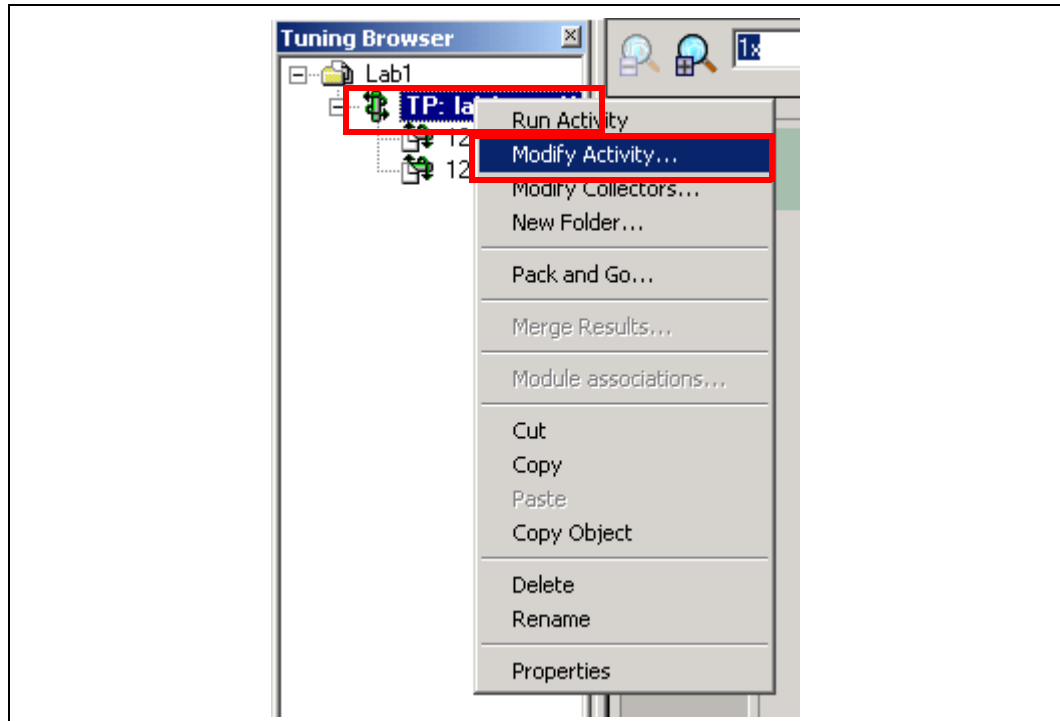
**Figure B-7. Specifying Timeline Information Collection Level**



## Creating a New Activity by Modifying an Existing Activity

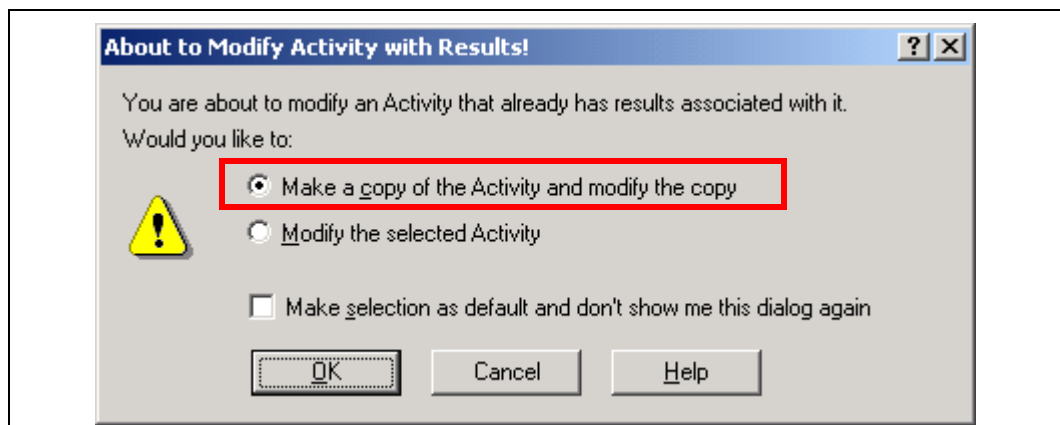
1. Right click on the current activity in the "Tuning Browser" window and select "Modify Activity". (Figure B-8)

**Figure B-8. Modifying an Existing Activity**



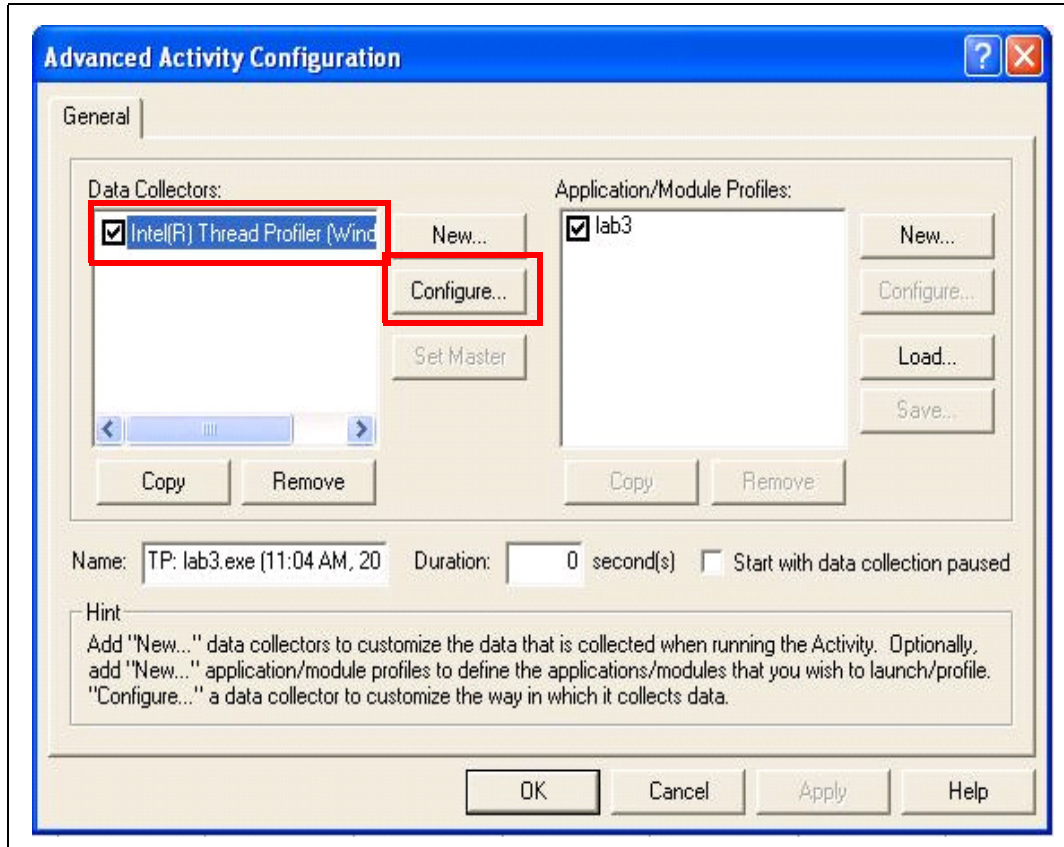
2. The About to Modify Activity pop-up dialog appears as shown in Figure B-9. Ensure that the "Make a copy of the Activity and modify the copy" radio-button is selected.

**Figure B-9. About to Modify Activity dialog**



3. Select "Thread Profiler" and click on "Configure" (Figure B-10). In the resulting pop-up dialog, select the "Miscellaneous" Tab.

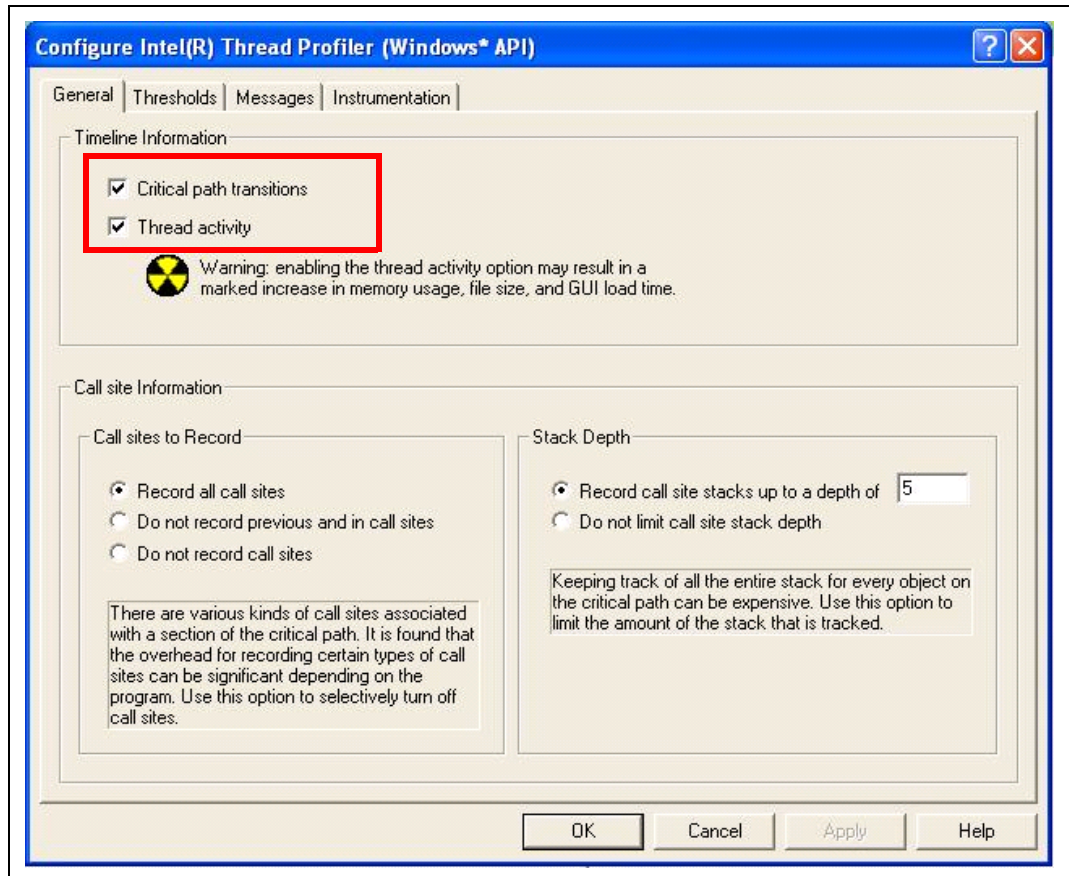
**Figure B-10. Advanced Activity Configuration**





4. Select the "Thread Activity" and "Transitions" check boxes (Figure B-11). This selection will issue a warning that the Thread Profiler may take longer to run. Accept this warning by clicking on "Yes".

**Figure B-11. Specifying Timeline Information Collection Level**





## Additional Resources

**new**

### ***Introduction to Parallel Programming***

Starting from foundation principles, this course introduces concepts and approaches common to all implementations of parallel programming for shared-memory systems.


Topics include Recognizing parallelism opportunities, dealing with sequential constructs, using threads to implement data and functional parallelism, discovering dependencies and ensuring mutual exclusion, analyzing and improving threaded performance, and choosing an appropriate threading model for implementation.

- For more information on Intel Software College, visit [www.intel.com/software/college](http://www.intel.com/software/college) .
- For more information on software development products, services, tools, training and expert advice, visit [www.intel.com/software](http://www.intel.com/software) .
- Put your knowledge to the ultimate test solving coding problems in competitions for multi-threading on multi-core microprocessors and win cash prizes! Intel Multi-Threading Competition Series at [www.topcoder.com/intel](http://www.topcoder.com/intel) .
- For more information about the latest technologies for computer product developers and IT professionals, look up Intel Press books at <http://www.intel.com/intelpress/> .
- Maximize application performance using Intel® Software Development Products: [www.intel.com/software/products/](http://www.intel.com/software/products/) .









[www.intel.com/software/college](http://www.intel.com/software/college)

Copyright © 2006, Intel Corporation. All rights reserved.

□ Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries.

\*Other brands and names are the property of their respective owners. □ □